

Never miss a tutorial:



MACHIN  
E LEARN  
ING MASTERY  
f ch i M Learning Mastery  
Making Developers Awesome at Machine Learning

Picked for you:

 How to Develop a Pix2Pix GAN for Image-to-Image Translation

Click to Take the FREE GANs Crash-Course

Search...



How to Develop a 1D Generative Adversarial Network From Scratch in Keras

# How to Develop a Conditional GAN (cGAN) From Scratch

 How to Develop a CycleGAN for Image-to-Image Translation with Keras  
Brownlee on July 5, 2019 in Generative Adversarial Networks

Tweet

Tweet

Share

Share

Published September 1, 2020  
How to Develop a Conditional GAN (cGAN) From Scratch

Generative Adversarial Networks, or GANs, are an architecture for training generative models, such as deep convolutional neural networks for generating images.

How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch

GAN models are capable of generating new random plausible examples for a given dataset, there is no way to control the types of images that are generated other than trying to figure out the complex relationship between the latent space input to the generator and the generated images.

## Loving the Tutorials?

The conditional generative adversarial network, or cGAN for short, is a type of GAN that involves the conditional generation of images by a generator model. Image generation can be conditional on a class label, if available, allowing the targeted generation of images of a given type.

In this tutorial, we will learn how to develop a conditional generative adversarial network for the targeted generation of items of clothing.

After completing this tutorial, you will know:

- The limitations of generating random samples with a GAN that can be overcome with a conditional generative adversarial network.
- How to develop and evaluate an unconditional generative adversarial network for generating photos of items of clothing.
- How to develop and evaluate a conditional generative adversarial network for generating photos of items of clothing.

**Kick-start your project** with my new book **Generative Adversarial Networks with Python**, including step-by-step tutorials and the *Python source code files* for all examples.

Let's get started.

**Never miss a tutorial:**

Photo by Big Cypress National Preserve, some rights reserved



This tutorial is divided into five parts; they are:

- 1 How to Develop a Pix2Pix GAN
- 2 Fashion Image-to-Image Translation Dataset
- 3 Unconditional GAN for Fashion-MNIST
- 4 Conditional GAN for Fashion-MNIST
  - How to Develop a 1D Generative Adversarial Network From Scratch in Keras

## Conditional Generative Adversarial Networks

How to Develop a CycleGAN for Image-to-Image Translation with Keras

The architecture is comprised of a generator and a discriminator model. The generator model is responsible for generating new plausible examples that ideally are indistinguishable from real examples in the dataset. The discriminator model is responsible for classifying a given image as either real (drawn from the dataset) or fake (generated).

How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch

discriminator trained together in a zero-sum or adversarial manner, such that improvements in the discriminator From Scratch at the cost of a reduced capability of the generator, and vice versa.

GANs are effective at image synthesis, that is, generating new examples of images for a target dataset. Some data is available for training, such as a class label, and it is desirable to make use of this information.

The [GANs with Python EBook](#) is where you'll find the [Really Good](#) stuff. For example, the MNIST handwritten digit dataset has class labels of the corresponding integers, the CIFAR-10 dataset has class labels for the corresponding objects in the photographs, and the Fashion-MNIST clothing dataset has class labels for the corresponding items of clothing.

There are two motivations for making use of the class label information in a GAN model.

1. Improve the GAN.
2. Targeted Image Generation.

Additional information that is correlated with the input images, such as class labels, can be used to improve the GAN. This improvement may come in the form of more stable training, faster training, and/or generated images that have better quality.

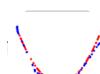
Class labels can also be used for the deliberate or targeted generation of images of a given type.

A limitation of a GAN model is that it may generate a random image from the domain. There is a relationship between points in the latent space to the generated images, but this relationship is complex and hard to map.

Alternately, a GAN can be trained in such a way that both the generator and the discriminator models are conditioned on the class label. This means that when the trained generator model is used as a module to generate images in the domain, images of a given type, or class label, can be generated.

### Picked for you:

 [Generative adversarial nets can be extended to a conditional model if both the generator and discriminator are conditioned on some extra information y. \[...\] We can perform the conditioning by feeding y into the both the discriminator and generator as additional input layer.](#)

 [How to Develop a 1D Generative Adversarial Networks, 2014.](#)

Keras

For example, in the case of MNIST, specific handwritten digits can be generated, such as the number 9; in the case of CIFAR-10, specific object photographs can be generated such as ‘frogs’; and in the case of the [Fashion MNIST dataset](#), specific items of clothing can be generated, such as ‘dress.’

This type of model is called a Conditional Generative Adversarial Network, CGAN or cGAN for short.

 [How to Develop a Conditional GAN](#)  
AN was first described by Mehdi Mirza and Simon Osindero in their 2014 paper titled “Conditional Generative Adversarial Nets.” In the paper, the authors motivate the approach based on the desire to direct the image generation process of the generator model.

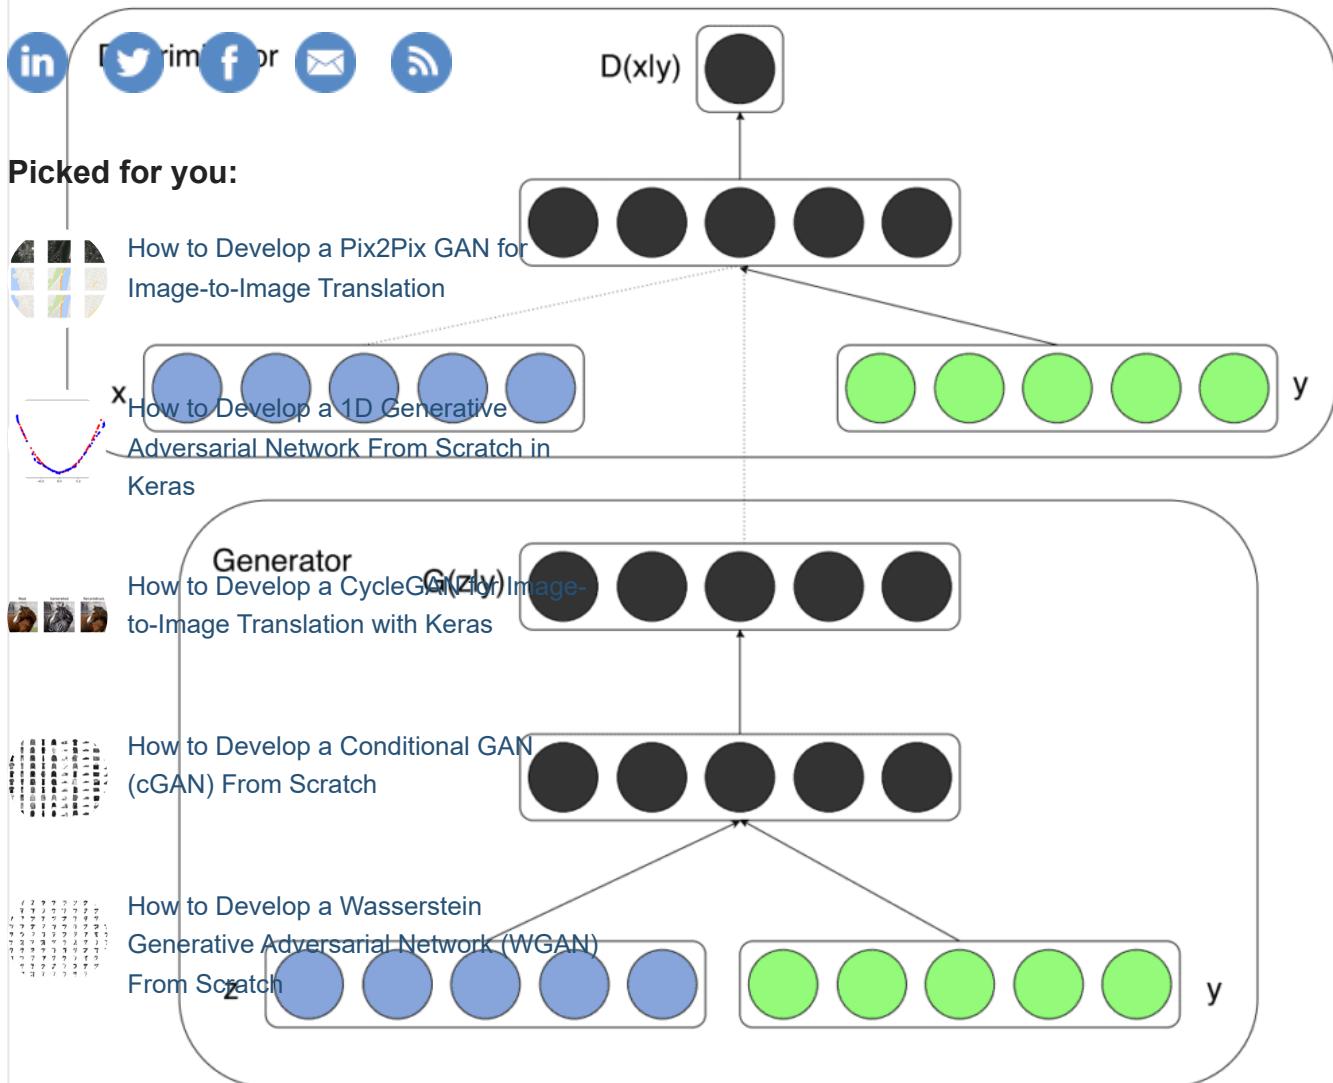
 [How to Develop a Wasserstein Generative Adversarial Network \(WGAN\)](#)  
... by conditioning the model on additional information it is possible to direct the data generation process. Such conditioning could be based on class labels

— [Conditional Generative Adversarial Nets, 2014.](#)  
**Loving the Tutorials?**

Their approach is demonstrated on the MNIST handwritten digit dataset where the class labels are one hot encoded and concatenated with the input to both the generator and discriminator models.

The image [>> SEE WHAT'S INSIDE](#) shows a diagram of the model architecture.

## Never miss a tutorial:



## Loving the Tutorials?

Example of a Conditional Generator and a Conditional Discriminator in a Conditional Generative Adversarial Network. The [GANs with Python EBook](#) is taken from [Conditional Generative Adversarial Nets](#), 2014, where you'll find the [Really Good](#) stuff.

There have been many improvements in the design and training of GAN models, most notably the [deep convolutional GAN](#), or DCGAN for short, that outlines the model configuration and training procedures that reliably result in the stable training of GAN models for a wide variety of problems. The conditional training of the DCGAN-based models may be referred to as CDCGAN or cDCGAN for short.

There are many ways to encode and incorporate the class labels into the discriminator and generator models. A best practice involves using an embedding layer followed by a fully connected layer with a linear activation that scales the embedding to the size of the image before concatenating it in the model as an additional channel or feature map.

A version of this recommendation was described in the 2015 paper titled “[Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks](#).”

“... we also explore a class conditional version of the model, where a vector  $c$  encodes the label. This is integrated into  $G_k$  &  $D_k$  by passing it through a linear layer whose output is

*reshaped into a single plane feature map which is then concatenated with the 1st layer maps.*  
**Never miss a tutorial:**



This recommendation was later added to the ‘[GAN Hacks](#)’ list of heuristic recommendations when **Picked for you:** designing and training GAN models, summarized as:

- [How to Develop a Pix2Pix GAN for 1D Discrete Variables in Conditional GANs](#)
  - Use an Embedding layer
  - Add as additional channels to images
  - [How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)
- [GAN Hacks](#)

 [How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)  
 In GANs can be conditioned on the class label, so-called class-conditional GANs, they can also be conditioned on other inputs, such as an image, in the case where a GAN is used for image-to-image translation tasks.

 [How to Develop a Conditional GAN \(cGAN\) From Scratch](#)  
 In this tutorial we will develop a GAN, specifically a DCGAN, then update it to use class labels in a conditional manner, specifically a cDCGAN model architecture.

 [How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)

## Want to Develop GANs from Scratch?

Take my free 7-day email crash course now (with sample code).

### Loving the Tutorials?

Click to sign-up and also get a free PDF Ebook version of the course.

The [GANs with Python](#) EBook is

where you'll find the **Really Good** stuff.

[Download Your FREE Mini-Course](#)

[>> SEE WHAT'S INSIDE](#)

## Fashion-MNIST Clothing Photograph Dataset

The [Fashion-MNIST](#) dataset is proposed as a more challenging replacement dataset for the MNIST dataset.

It is a dataset comprised of 60,000 small square 28×28 pixel grayscale images of items of 10 types of clothing, such as shoes, t-shirts, dresses, and more.

Keras provides access to the Fashion-MNIST dataset via the `fashion_mnist.load_dataset()` function. It returns two tuples, one with the input and output elements for the standard training dataset, and another with the input and output elements for the standard test dataset.

The example below loads the dataset and summarizes the shape of the loaded dataset.

Note: the first time you load the dataset, Keras will automatically download a compressed version of the images and save them under your home directory in `~/.keras/datasets/`. The download is fast as the dataset is only 2 megabytes in its compressed form.



```

1 # example of loading the fashion_mnist dataset
2 from keras.datasets.fashion_mnist import load_data
3 # Load the images into memory
4 (trainX, trainy), (testX, testy) = load_data()
5 # summarize the shape of the dataset
6 print('Train', trainX.shape, trainy.shape)
7 print('Test', testX.shape, testy.shape)

```

Running the example loads the dataset and prints the shape of the input and output components of the

[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

We can see that there are 60K examples in the training set and 10K in the test set and that each image is a square of 28 by 28 pixels.

[How to Develop a CycleGAN for Image-to-Image Translation](#)

```

1 Train (60000, 28, 28) (60000, 10)
2 Test (10000, 28, 28) (10000, 10)

```

The images are grayscale with a black background (0 pixel value) and the items of clothing are in white

[How to Develop a Conditional GAN \(cGAN\) From Scratch](#) (values near 255). This means if the images were plotted, they would be mostly black with a white clothing in the middle.

We can plot some of the images from the training dataset using the `matplotlib` library with the `imshow()`

[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)

[From Scratch](#)

```

1 # plot raw pixel data
2 pyplot.imshow(trainX[i], cmap='gray')

```

Alternately, the images are easier to review when we reverse the colors and plot the background as white and the clothing in black.

[The GANs with Python EBook](#)

[They are easier to view because the image is now white with the area of interest in black. This can be achieved using a reverse grayscale color map, as follows:](#)

[>> SEE WHAT'S INSIDE](#)

```

1 # plot raw pixel data
2 pyplot.imshow(trainX[i], cmap='gray_r')

```

The example below plots the first 100 images from the training dataset in a 10 by 10 square.

```

1 # example of loading the fashion_mnist dataset
2 from keras.datasets.fashion_mnist import load_data
3 from matplotlib import pyplot
4 # load the images into memory
5 (trainX, trainy), (testX, testy) = load_data()
6 # plot images from the training dataset
7 for i in range(100):
8     # define subplot
9     pyplot.subplot(10, 10, 1 + i)
10    # turn off axis
11    pyplot.axis('off')
12    # plot raw pixel data
13    pyplot.imshow(trainX[i], cmap='gray_r')
14    pyplot.show()

```

Running the example creates a figure with a plot of 100 images from the MNIST training dataset, **Never miss a tutorial!** arranged in a  $10 \times 10$  square.



### Picked for you:



### Loving the Tutorials?

Plot of the First 100 Items of Clothing From the Fashion MNIST Dataset.

The [GANs with Python](#) EBook is  
We will use the images in the training dataset as the basis for training a Generative Adversarial Network.  
~~Where you'll make really good stuff.~~

[>> SEE WHAT'S INSIDE](#)

Specifically, the generator model will learn how to generate new plausible items of clothing using a discriminator that will try to distinguish between real images from the Fashion MNIST training dataset and new images output by the generator model.

This is a relatively simple problem that does not require a sophisticated generator or discriminator model, although it does require the generation of a grayscale output image.

## Unconditional GAN for Fashion-MNIST

In this section, we will develop an unconditional GAN for the Fashion-MNIST dataset.

The first step is to define the models.

The discriminator model takes as input one  $28 \times 28$  grayscale image and outputs a binary prediction as to whether the image is real (class=1) or fake (class=0). It is implemented as a modest convolutional neural network using best practices for GAN design such as using the LeakyReLU activation function

with a slope of 0.2, using a  $2 \times 2$  stride to downsample, and the adam version of stochastic gradient descent with a learning rate of 0.0002 and a momentum of 0.5



The `def define_discriminator()` function below implements this, defining and compiling the discriminator model and returning it. The input shape of the image is parameterized as a default function argument in **Picked for you**: re-use the function for your own image data later.

```

1 # How to Develop a Pix2Pix GAN for
2 def define_discriminator(in_shape=(28,28,1)):
3     model = Sequential()
4     # downsample
5     model.add(Conv2D(128, (3,3), strides=(2,2), padding='same', input_shape=in_shape))
6     model.add(LeakyReLU(alpha=0.2))
7     # downsample
8     model.add(Conv2D(128, (3,3), strides=(2,2), padding='same'))
9     model.add(LeakyReLU(alpha=0.2))
10    # classifier
11    model.add(Flatten())
12    model.add(Dropout(0.4))
13    model.add(Dense(1, activation='sigmoid'))
14    # compile model
15    opt = Adam(lr=0.0002, beta_1=0.5)
16    model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
17    return model

```

The generator model takes as input a point in the latent space and outputs a single  $28 \times 28$  grayscale image. This is achieved by using a fully connected layer to interpret the point in the latent space and provide sufficient activations that can be reshaped into many copies (in this case 128) of a low-resolution version of the output image (e.g.  $7 \times 7$ ). This is then upsampled twice, doubling the size and quadrupling the area of the activations each time using transpose convolutional layers. The model uses best practices such as the LeakyReLU activation, a kernel size that is a factor of the stride size, and a hyperbolic tangent (tanh) activation function in the output layer.

### Loving the Tutorials?

The `define_generator()` function below defines the generator model, but intentionally does not compile it as it is not trained directly; then returns the model. The size of the latent space is parameterized as a function argument.

```

1 # define >> SEE WHAT'S INSIDE << the standalone generator model
2 def define_generator(latent_dim):
3     model = Sequential()
4     # foundation for 7x7 image
5     n_nodes = 128 * 7 * 7
6     model.add(Dense(n_nodes, input_dim=latent_dim))
7     model.add(LeakyReLU(alpha=0.2))
8     model.add(Reshape((7, 7, 128)))
9     # upsample to 14x14
10    model.add(Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
11    model.add(LeakyReLU(alpha=0.2))
12    # upsample to 28x28
13    model.add(Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
14    model.add(LeakyReLU(alpha=0.2))
15    # generate
16    model.add(Conv2D(1, (7,7), activation='tanh', padding='same'))
17    return model

```

Next, a GAN model can be defined that combines both the generator model and the discriminator model into one larger model. This larger model will be used to train the model weights in the generator, using the output and error calculated by the discriminator model. The discriminator model is trained separately, and as such, the model weights are marked as not trainable in this larger GAN model to

ensure that only the weights of the generator model are updated. This change to the trainability of the discriminator weights only has an effect when training the combined GAN model, not when training the generator model.



This larger GAN model takes as input a point in the latent space, uses the generator model to generate images, which are then fed as input to the discriminator model, then is output or classified as real or fake.

The `define_gan()` function below implements this, taking the already-defined generator and discriminator models as input.

```

1 # define the combined generator and discriminator model, for updating the generator
2 def define_gan(generator, discriminator):
3     # make weights in the discriminator not trainable
4     discriminator.trainable = False
5     # connect them
6     model = Sequential()
7     # add generator
8     model.add(generator)
9     # add the discriminator
10    model.add(discriminator)
11    # compile model
12    opt = Adam(lr=0.0002, beta_1=0.5)
13    model.compile(loss='binary_crossentropy', optimizer=opt)
14    return model

```

Now that we have defined the GAN model, we need to train it. But, before we can train the model, we require input data.

#### How to Develop a Wasserstein Generative Adversarial Network (WGAN)

The first step is to load and prepare the Fashion MNIST dataset. We only require the images in the training dataset. The images are black and white, therefore we must add an additional channel dimension to transform them to be three dimensional, as expected by the convolutional layers of our models. Finally, the pixel values must be scaled to the range [-1,1] to match the output of the generator model.

#### Loving the Tutorials?

The `GANs with Python` EBook is

The `load_real_samples()` function below implements this, returning the loaded and scaled Fashion MNIST training dataset ready for modeling.

>> SEE WHAT'S INSIDE

```

1 # load fashion mnist images
2 def load_real_samples():
3     # load dataset
4     (trainX, _), (_, _) = load_data()
5     # expand to 3d, e.g. add channels
6     X = expand_dims(trainX, axis=-1)
7     # convert from ints to floats
8     X = X.astype('float32')
9     # scale from [0,255] to [-1,1]
10    X = (X - 127.5) / 127.5
11    return X

```

We will require one batch (or a half) batch of real images from the dataset each update to the GAN model. A simple way to achieve this is to select a random sample of images from the dataset each time.

The `generate_real_samples()` function below implements this, taking the prepared dataset as an argument, selecting and returning a random sample of Fashion MNIST images and their corresponding class label for the discriminator, specifically class=1, indicating that they are real images.

```

1 # select real samples

```

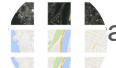
```

2 def generate_real_samples(dataset, n_samples):
3     # choose random instances
4     ix = randint(0, dataset.shape[0], n_samples)
5     # select images
6     X = dataset[ix]
7     # generate class labels
8     y = ones((n_samples, 1))
9     return X, y

```

locked for you:

Next, we need inputs for the generator model. These are random points from the latent space,

 How to Develop a Pix2Pix GAN for  
Ganally Gaussian distributed random variables.  
 Image-to-Image Translation

The `generate_latent_points()` function implements this, taking the size of the latent space as an argument and the number of points required and returning them as a batch of input samples for the generator model.

 Adversarial Network From Scratch in

Keras

```

1 # generate points in latent space as input for the generator
2 def generate_latent_points(latent_dim, n_samples):
3     # generate points in the latent space
4     x_input = randn(latent_dim * n_samples)
5     # reshape into a batch of inputs for the network
6     x_input = x_input.reshape(n_samples, latent_dim)
7     return x_input

```

 How to Develop a Conditional GAN  
We need to use the points in the latent space as input to the generator in order to generate new  
(cGAN) From Scratch

The `generate_fake_samples()` function below implements this, taking the generator model and size of the latent space as arguments, then generating points in the latent space and using them as input to the generator model. The function returns the generated images and their corresponding class label for the discriminator model, specifically class=0 to indicate they are fake or generated.

```

1 # use the generator to generate n fake examples, with class labels
2 def generate_fake_samples(generator, latent_dim, n_samples):
3     # generate points in latent space
4     x_input = generate_latent_points(latent_dim, n_samples)
5     # predict outputs
6     X = generator.predict(x_input)
7     # create class labels
8     y = zeros((n_samples, 1))
9     return X, y

```

We are now ready to fit the GAN models.

The model is fit for 100 training epochs, which is arbitrary, as the model begins generating plausible items of clothing after perhaps 20 epochs. A batch size of 128 samples is used, and each training epoch involves 60,000/128, or about 468 batches of real and fake samples and updates to the model.

First, the discriminator model is updated for a half batch of real samples, then a half batch of fake samples, together forming one batch of weight updates. The generator is then updated via the composite gan model. Importantly, the class label is set to 1 or real for the fake samples. This has the effect of updating the generator toward getting better at generating real samples on the next batch.

The `train()` function below implements this, taking the defined models, dataset, and size of the latent dimension as arguments and parameterizing the number of epochs and batch size with default arguments. The generator model is saved at the end of training.

```

1 # train the generator and discriminator
2 def train(g_model, d_model, gan_model, dataset, latent_dim, n_epochs=100, n_batch=128):
3     bat_per_epo = int(dataset.shape[0] / n_batch)
4     half_batch = int(n_batch / 2)
5     # manually enumerate epochs
6     for i in range(n_epochs):
7         # enumerate batches over the training set
8         for j in range(bat_per_epo):
9             # get randomly selected 'real' samples
10            X_real, y_real = generate_real_samples(dataset, half_batch)
11            # update discriminator model weights
12            d_loss1, _ = d_model.train_on_batch(X_real, y_real)
13            # generate 'fake' examples
14            X_fake, y_fake = generate_fake_samples(g_model, latent_dim, half_batch)
15            # update discriminator model weights
16            d_loss2, _ = d_model.train_on_batch(X_fake, y_fake)
17            # prepare points in latent space as input for the generator
18            X_gan = generate_latent_points(latent_dim, n_batch)
19            # create inverted labels for the fake samples
20            y_gan = ones((n_batch, 1))
21            # update the generator via the discriminator's error
22            g_loss = gan_model.train_on_batch(X_gan, y_gan)
23            # summarize loss on this batch
24            print('>%d, %d/%d, d1=%f, d2=%f g=%f' %
25                  (i+1, j+1, bat_per_epo, d_loss1, d_loss2, g_loss))
26            # save the generator model
27            g_model.save('generator.h5')

```

Then define the size of the latent space, define all three models, and train them on the loaded fashion MNIST dataset.

```

1 # size of the latent space
2 latent_dim = 100
3 # create the discriminator
4 discriminator = define_discriminator()
5 # create the generator
6 generator = define_generator(latent_dim)
7 # create the gan
8 gan_model = define_gan(generator, discriminator)
9 # load image data
10 dataset = load_real_samples()
11 # train model
12 train(generator, discriminator, gan_model, dataset, latent_dim)

```

Tying all this together, the example is listed below.

```

1 # example of training an unconditional gan on the fashion mnist dataset
2 from numpy import expand_dims
3 from numpy import zeros
4 from numpy import ones
5 from numpy.random import randn
6 from numpy.random import randint
7 from keras.datasets.fashion_mnist import load_data
8 from keras.optimizers import Adam
9 from keras.models import Sequential
10 from keras.layers import Dense
11 from keras.layers import Reshape
12 from keras.layers import Flatten
13 from keras.layers import Conv2D
14 from keras.layers import Conv2DTranspose
15 from keras.layers import LeakyReLU
16 from keras.layers import Dropout
17
18 # define the standalone discriminator model
19 def define_discriminator(in_shape=(28,28,1)):
20     model = Sequential()
21     # downsample

```

```

22 model.add(Conv2D(128, (3,3), strides=(2,2), padding='same', input_shape=in_shape))
23 model.add(LeakyReLU(alpha=0.2))
24 # downsample
25 model.add(Conv2D(128, (3,3), strides=(2,2), padding='same'))
26 model.add(LeakyReLU(alpha=0.2))
27 # classifier
28 model.add(Flatten())
29 model.add(Dropout(0.4))
30 model.add(Dense(1, activation='sigmoid'))
31 # compile model
32 opt = Adam(lr=0.0002, beta_1=0.5)
33 model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
34 return model
35
36 # define the standalone generator model
37 def define_generator(latent_dim):
38     model = Sequential()
39     # foundation for 7x7 image
40     n_nodes = 128 * 7 * 7
41     model.add(Dense(n_nodes, input_dim=latent_dim))
42     model.add(LeakyReLU(alpha=0.2))
43     model.add(Reshape((7, 7, 128)))
44     # upsample to 14x14
45     model.add(Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
46     model.add(LeakyReLU(alpha=0.2))
47     # upsample to 28x28
48     model.add(Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     # generate
51     model.add(Conv2D(1, (7,7), activation='tanh', padding='same'))
52     return model
53
54 # define the combined generator and discriminator model, for updating the generator
55 def define_gan(generator, discriminator):
56     # make weights in the discriminator not trainable
57     discriminator.trainable = False
58     # connect them
59     model = Sequential()
60     # add generator
61     model.add(generator)
62     # add the discriminator
63     model.add(discriminator)
64     # compile model
65     opt = Adam(lr=0.0002, beta_1=0.5)
66     model.compile(loss='binary_crossentropy', optimizer=opt)
67     return model
68
69 # load fashion mnist images
70 def load_real_samples():
71     # load dataset
72     (trainX, _), (_, _) = load_data()
73     # expand to 3d, e.g. add channels
74     X = expand_dims(trainX, axis=-1)
75     # convert from ints to floats
76     X = X.astype('float32')
77     # scale from [0,255] to [-1,1]
78     X = (X - 127.5) / 127.5
79     return X
80
81 # select real samples
82 def generate_real_samples(dataset, n_samples):
83     # choose random instances
84     ix = randint(0, dataset.shape[0], n_samples)
85     # select images
86     X = dataset[ix]
87     # generate class labels
88     y = ones((n_samples, 1))
89     return X, y
90

```

```

91 # generate points in latent space as input for the generator
92 def generate_latent_points(latent_dim, n_samples):
93     # generate points in the latent space
94     x_input = randn(latent_dim * n_samples)
95     # reshape into a batch of inputs for the network
96     x_input = x_input.reshape(n_samples, latent_dim)
97     return x_input
98
99 # use the generator to generate n fake examples, with class labels
100 def generate_fake_samples(generator, latent_dim, n_samples):
101     # generate points in latent space
102     x_input = generate_latent_points(latent_dim, n_samples)
103     # predict outputs
104     X = generator.predict(x_input)
105     # create class labels
106     y = zeros((n_samples, 1))
107     return X, y
108
109 # train the generator and discriminator
110 def train(g_model, d_model, gan_model, dataset, latent_dim, n_epochs=100, n_batch=128):
111     bat_per_epo = int(dataset.shape[0] / n_batch)
112     half_batch = int(n_batch / 2)
113     # manually enumerate epochs
114     for i in range(n_epochs):
115         # enumerate batches over the training set
116         for j in range(bat_per_epo):
117             # get randomly selected 'real' samples
118             X_real, y_real = generate_real_samples(dataset, half_batch)
119             # update discriminator model weights
120             d_loss1, _ = d_model.train_on_batch(X_real, y_real)
121             # generate 'fake' examples
122             X_fake, y_fake = generate_fake_samples(g_model, latent_dim, half_batch)
123             # update discriminator model weights
124             d_loss2, _ = d_model.train_on_batch(X_fake, y_fake)
125             # prepare points in latent space as input for the generator
126             X_gan = generate_latent_points(latent_dim, n_batch)
127             # create inverted labels for the fake samples
128             y_gan = ones((n_batch, 1))
129             # update the generator via the discriminator's error
130             g_loss = gan_model.train_on_batch(X_gan, y_gan)
131             # summarize loss on this batch
132             print('>%d, %d/%d, d1=%f, d2=%f, g=%f' %
133                   (i+1, j+1, bat_per_epo, d_loss1, d_loss2, g_loss))
134             # save the generator model
135             g_model.save('generator.h5')
136
137     # size of the latent space
138     latent_dim = 100
139     # create the discriminator
140     discriminator = define_discriminator()
141     # create the generator
142     generator = define_generator(latent_dim)
143     # create the gan
144     gan_model = define_gan(generator, discriminator)
145     # load image data
146     dataset = load_real_samples()
147     # train model
148     train(generator, discriminator, gan_model, dataset, latent_dim)

```

Running the example may take a long time on modest hardware.

I recommend running the example on GPU hardware. If you need help, you can get started quickly by using an AWS EC2 instance to train the model. See the tutorial:

- How to Setup Amazon AWS EC2 GPUs to Train Keras Deep Learning Models (step-by-step)

The loss for the discriminator on real and fake samples, as well as the loss for the generator, is reported after each batch.



Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average

## Picked for you:

How to Develop a Pix2Pix GAN for Image-to-Image Translation with Keras

```

1 ...
2 >100, 464/468, d1=0.681, d2=0.685 g=0.693
3 >100, 465/468, d1=0.691, d2=0.700 g=0.703
4 >100, 466/468, d1=0.691, d2=0.703 g=0.706
5 >100, 467/468, d1=0.698, d2=0.699 g=0.699
6 >100, 468/468, d1=0.699, d2=0.695 g=0.708

```

How to Develop a CycleGAN for Image-to-Image Translation with Keras

This model can be loaded and used to generate new random but plausible samples from the fashion MNIST dataset.

How to Develop a Conditional GAN

(cGAN) From Scratch

The sample below loads the saved model and generates 100 random items of clothing.

```

1 # example of loading the generator model and generating images
2 from keras.models import load_model
3 from numpy.random import randn
4 from matplotlib import pyplot
5
6 # generate points in latent space as input for the generator
7 def generate_latent_points(latent_dim, n_samples):
8     # generate points in the latent space
9     x_input = randn(latent_dim * n_samples)
10    # reshape into a batch of inputs for the network
11    x_input = x_input.reshape(n_samples, latent_dim)
12    return x_input
13
14 # create and save a plot of generated images (reversed grayscale)
15 def show_plot(examples, n):
16     # plot images
17     for i in range(n * n):
18         # define subplot
19         pyplot.subplot(n, n, 1 + i)
20         # turn off axis
21         pyplot.axis('off')
22         # plot raw pixel data
23         pyplot.imshow(examples[i, :, :, 0], cmap='gray_r')
24     pyplot.show()
25
26 # load model
27 model = load_model('generator.h5')
28 # generate images
29 latent_points = generate_latent_points(100, 100)
30 # generate images
31 X = model.predict(latent_points)
32 # plot the result
33 show_plot(X, 10)

```

Running the example creates a plot of 100 randomly generated items of clothing arranged into a  $10 \times 10$  grid.

**Note:** Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average



In this case, we can see an assortment of clothing items such as shoes, sweaters, and pants. Most **Picked for you** plausible and could have come from the fashion MNIST dataset. They are not perfect, however, as there are some sweaters with a single sleeve and shoes that look like a mess.

[How to Develop a Pix2Pix GAN for](#)



[Image-to-Image Translation](#)

	<a href="#">How to Develop a 1D Generative Adversarial Network From Scratch in Keras</a>	
	<a href="#">How to Develop a CycleGAN for Image-to-Image Translation with Keras</a>	
	<a href="#">How to Develop a Conditional GAN (cGAN) From Scratch</a>	
	<a href="#">How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch</a>	
	<b>Loving the tutorials?</b>	
<a href="#">The GANs with Python EBook</a> is where you'll find the <b>Really Good</b> stuff.		

[>> SEE WHAT'S INSIDE](#)

Example of 100 Generated items of Clothing using an Unconditional GAN.

## Conditional GAN for Fashion-MNIST

In this section, we will develop a conditional GAN for the Fashion-MNIST dataset by updating the unconditional GAN developed in the previous section.

The best way to design models in Keras to have multiple inputs is by using the [Functional API](#), as opposed to the Sequential API used in the previous section. We will use the functional API to re-implement the discriminator, generator, and the composite model.

Starting with the discriminator model, a new second input is defined that takes an integer for the class label of the image. This has the effect of making the input image conditional on the provided class label.

The class label is then passed through an Embedding layer with the size of 50. This means that each of the 10 classes for the Fashion MNIST dataset (0 through 9) will map to a different 50-element vector

representation that will be learned by the discriminator model.

### Never miss a tutorial:

The output of the embedding is then passed to a fully connected layer with a linear activation.

Importantly, the fully connected layer has enough activations that can be reshaped into one channel of a 28×28 image. The activations are reshaped into single 28×28 activation map and concatenated with the image. This has the effect of looking like a two-channel input image to the next convolutional layer.

 How to Develop a Pix2Pix GAN for

 `image_discriminator()` below implements this update to the discriminator model. The parameterized shape of the input image is also used after the embedding layer to define the number of activations for the fully connected layer to reshape its output. The number of classes in the problem is also parameterized in the function and set

Adversarial Network From Scratch in

Keras

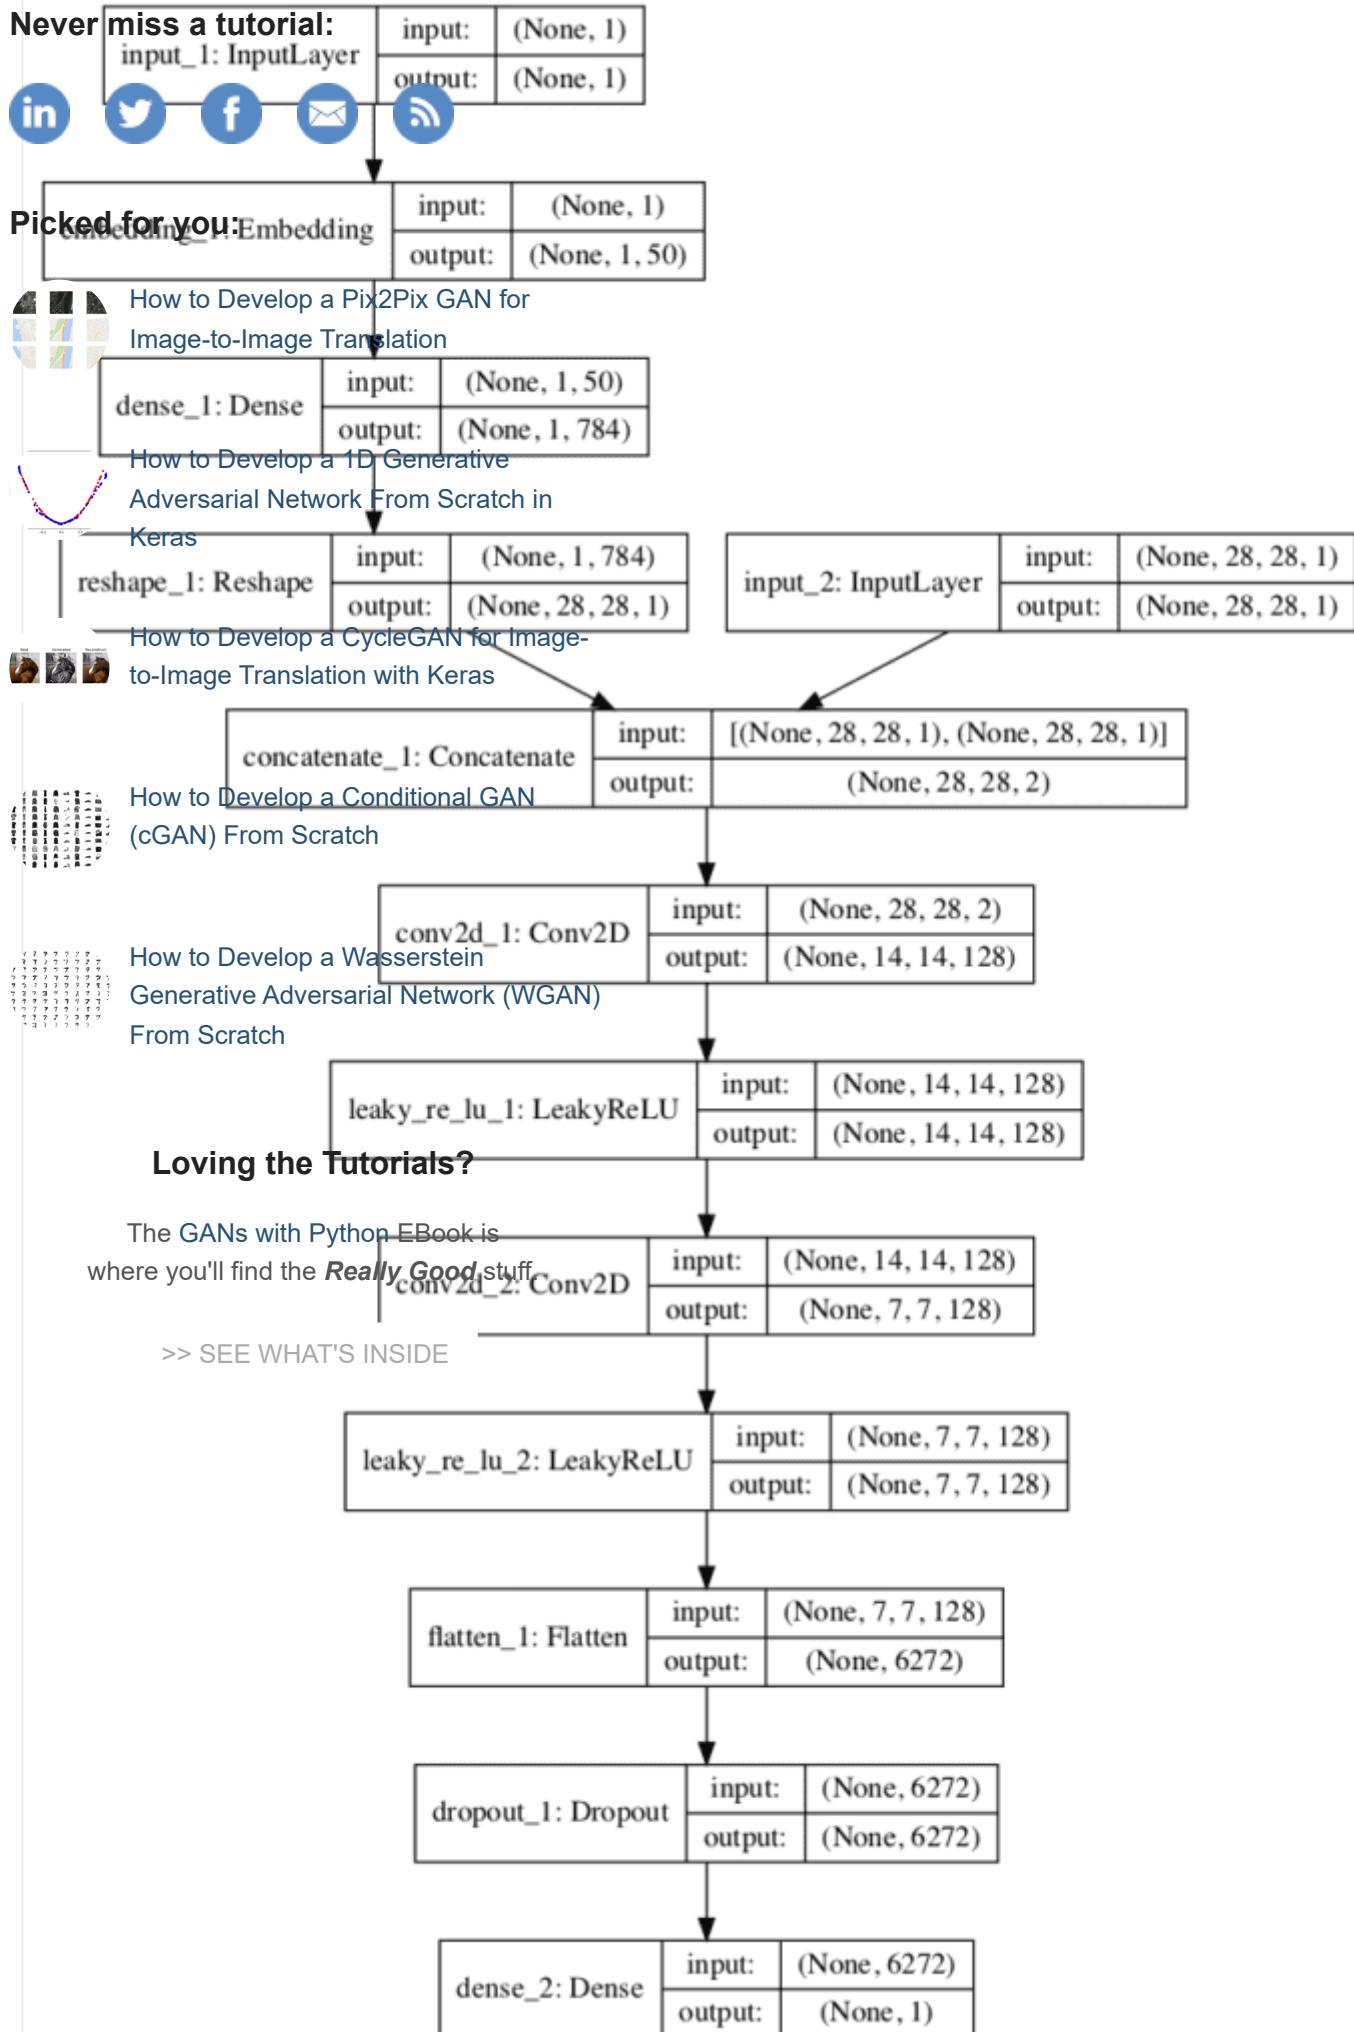
```

1 # define the standalone discriminator model
2 def define_discriminator(in_shape=(28,28,1), n_classes=10):
3     # label input
4     in_label = Input(shape=(1,))
5     # embedding for categorical input
6     li = Embedding(n_classes, 50)(in_label)
7     # scale up to image dimensions with linear activation
8     n_nodes = in_shape[0] * in_shape[1]
9     li = Dense(n_nodes)(li) # How to Develop a Conditional GAN
10    # reshape to additional channel
11    li = Reshape((in_shape[0], in_shape[1], 1))(li)
12    # image input
13    in_image = Input(shape=in_shape)
14    # concat label as a channel
15    merge = Concatenate()([in_image, li])
16    # downsample
17    fe = Conv2D(128, (3,3), strides=(2,2), padding='same')(merge)
18    fe = LeakyReLU(alpha=0.2)(fe)
19    # downsample
20    fe = Conv2D(128, (3,3), strides=(2,2), padding='same')(fe)
21    fe = LeakyReLU(alpha=0.2)(fe) # Leaving the Tutorial?
22    # flatten feature maps
23    fe = Flatten()(fe) # This is where you can learn Python F-Book is
24    # dropout
25    fe = Dropout(0.4)(fe) # where you'll type many many good stuff.
26    # output
27    out_layer = Dense(1, activation='sigmoid')(fe)
28    # define model
29    model = Model([in_image, in_label], out_layer)
30    # compile model
31    opt = Adam(lr=0.0002, beta_1=0.5)
32    model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
33    return model

```

In order to make the architecture clear, below is a plot of the discriminator model.

The plot shows the two inputs: first the class label that passes through the embedding (left) and the image (right), and their concatenation into a two-channel 28×28 image or feature map (middle). The rest of the model is the same as the discriminator designed in the previous section.



Plot of the Discriminator Model in the Conditional Generative Adversarial Network

**Never miss a tutorial:**

Next, the generator model must be updated to take the class label. This has the effect of making the point in latent space conditional on the provided class label.

**Picked for you:** As in the discriminator, the class label is passed through an embedding layer to map it to a unique 50-element vector and is then passed through a fully connected layer with a linear activation before being resized into a single  $7 \times 7$  feature map. This is to match the  $7 \times 7$  feature map activations of the unconditional generator model. The new  $7 \times 7$  feature map is added as one more channel to the existing 128, resulting in 129 feature maps that are then upsampled as in the prior model.

[How to Develop a 1D Generative](#)

[Adversarial Network From Scratch in Keras](#) function below implements this, again parameterizing the number of classes as we did with the discriminator model.

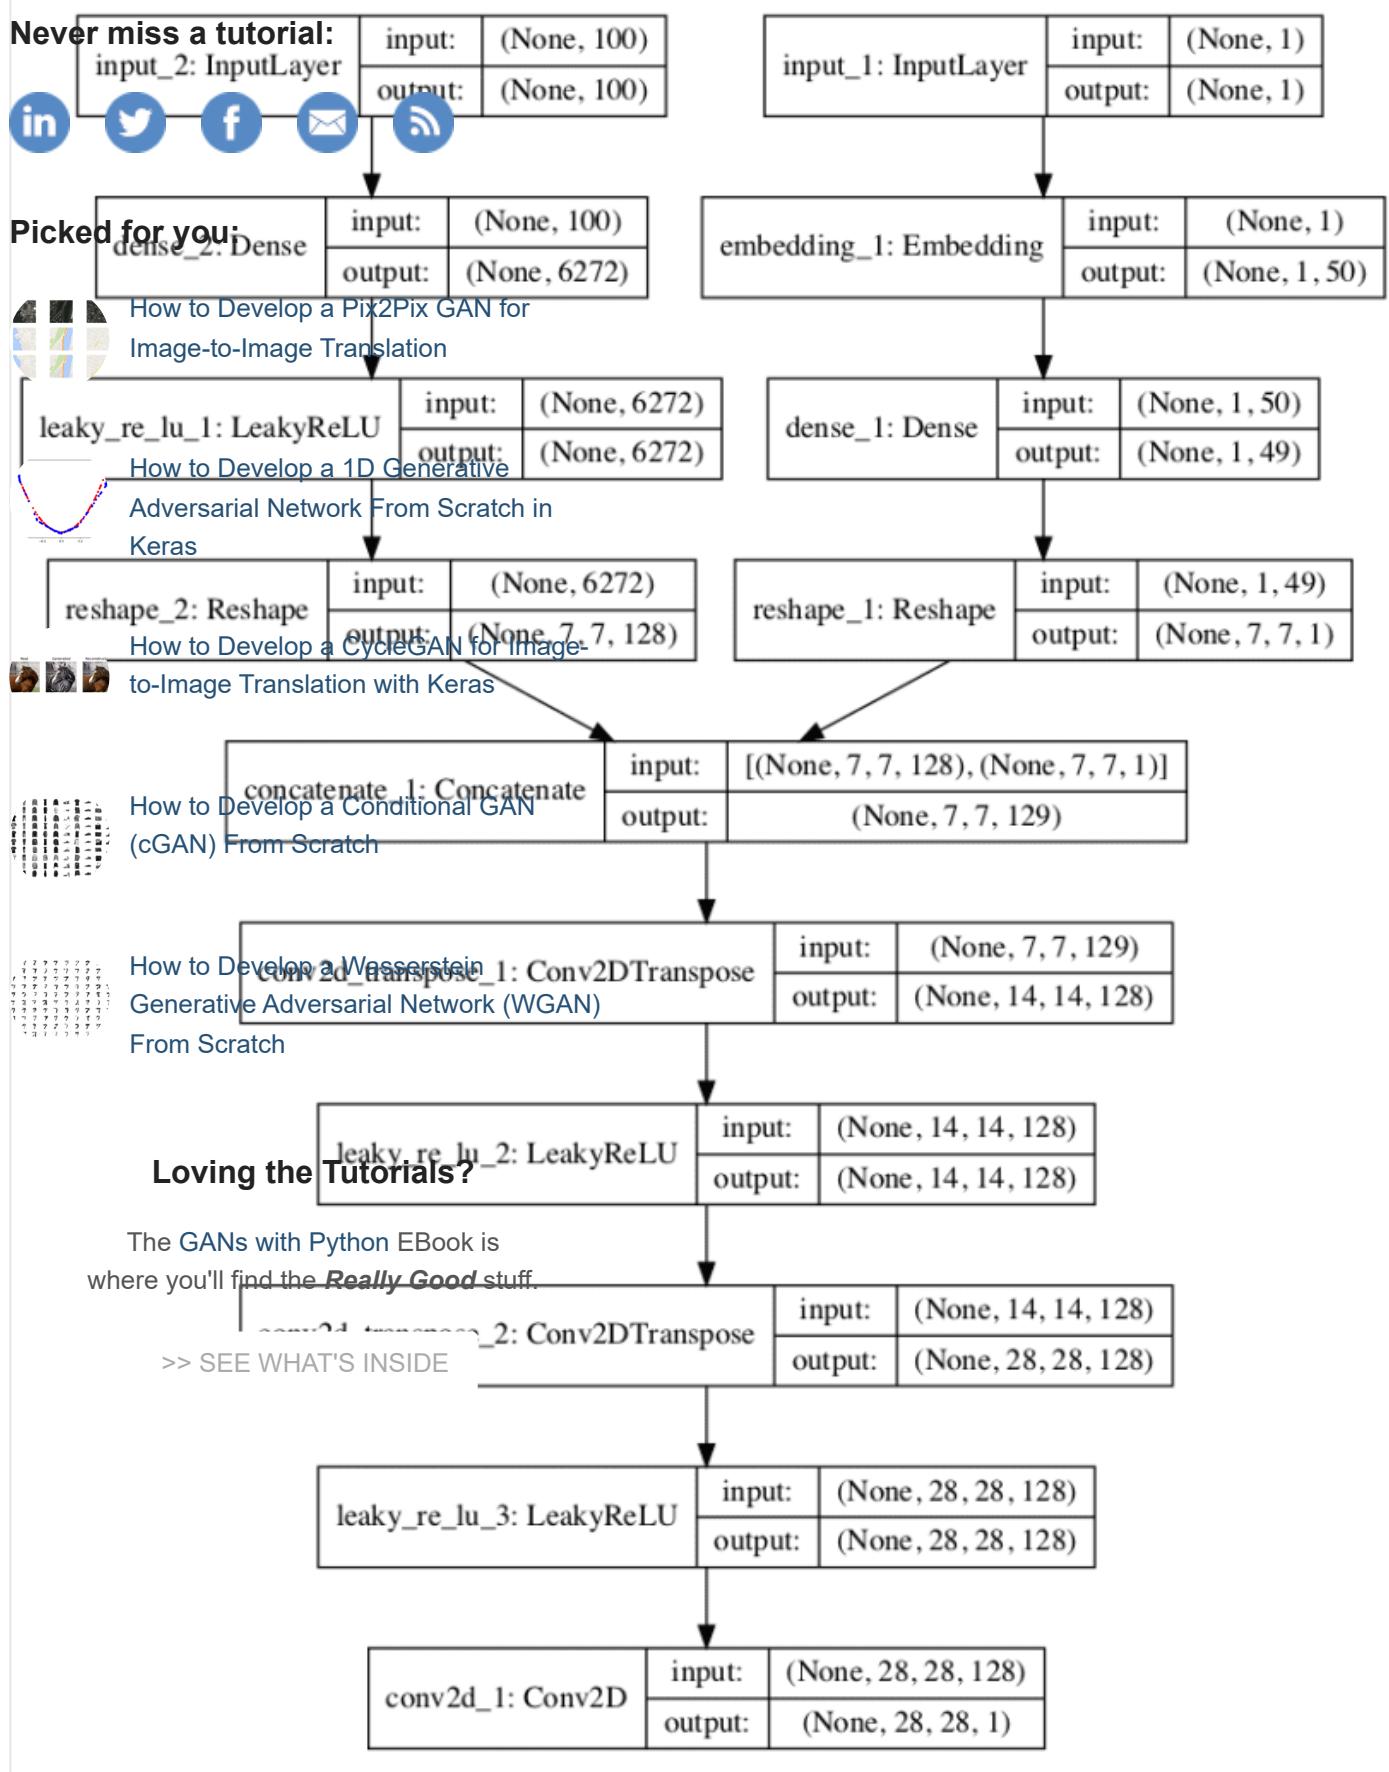
```

1 # How to Develop a CycleGAN for Image Model
2 def define_generator(latent_dim, n_classes=10):
3     # label input
4     in_label = Input(shape=(1,))
5     # embedding for categorical input
6     li = Embedding(n_classes, 50)(in_label)
7     # linear multiplication
8     n_nodes = 7 * 7
9     li = Dense(n_nodes)(li)
10    # reshape to additional channel
11    li = Reshape((7, 7, 1))(li)
12    # image generator input
13    in_Generator_Adversarial_Network(cGAN)
14    # foundation for 7x7 image
15    n_nodes = 128 * 7 * 7
16    gen = Dense(n_nodes)(in_lat)
17    gen = LeakyReLU(alpha=0.2)(gen)
18    gen = Reshape((7, 7, 128))(gen)
19    # merge image gen and label input
20    merge = Concatenate()([gen, li])
21    # upsample to 14x14
22    gen = Conv2DTranspose(128, (4,4), strides=(2,2), padding='same')(merge)
23    gen = LeakyReLU(alpha=0.2)(gen)
24    # upsample to 28x28
25    gen = Conv2DTranspose(128, (4,4), strides=(2,2), padding='same')(gen)
26    gen = LeakyReLU(alpha=0.2)(gen)
27    # output
28    out_layer = Conv2D(1, (7,7), activation='tanh', padding='same')(gen)
29    # define model
30    model = Model([in_lat, in_label], out_layer)
31    return model

```

To help understand the new model architecture, the image below provides a plot of the new conditional generator model.

In this case, you can see the 100-element point in latent space as input and subsequent resizing (left) and the new class label input and embedding layer (right), then the concatenation of the two sets of feature maps (center). The remainder of the model is the same as the unconditional case.



Finally, the composite GAN model requires updating.

The new GAN model will take a point in latent space as input and a class label and generate a prediction of whether input was real or fake, as before.

**Never miss a tutorial:** Using the functional API to design the model, it is important that we explicitly connect the image generated output from the generator as well as the class label input, both as input to the discriminator model. This allows the same class label input to flow down into the generator and down into the discriminator.

**Picked for you!** function below implements the conditional version of the GAN.

```

1 # Define the combined generator and discriminator model, for updating the generator
2 def define_gan(g_model, d_model):
3     # make weights in the discriminator not trainable
4     d_model.trainable = False
5     # get noise and label inputs from generator model
6     gen_noise, gen_label = g_model.input
7     # get image output from the generator model
8     gen_output = g_model.output
9     # connect image output and label input from generator as inputs to discriminator
10    gan_output = d_model([gen_output, gen_label])
11    # define gan model as taking noise and label and outputting a classification
12    model = Model([gen_noise, gen_label], gan_output)
13    # compile model
14    opt = Adam(lr=0.0002, beta_1=0.5)
15    model.compile(loss='binary_crossentropy', optimizer=opt)
16    return model

```

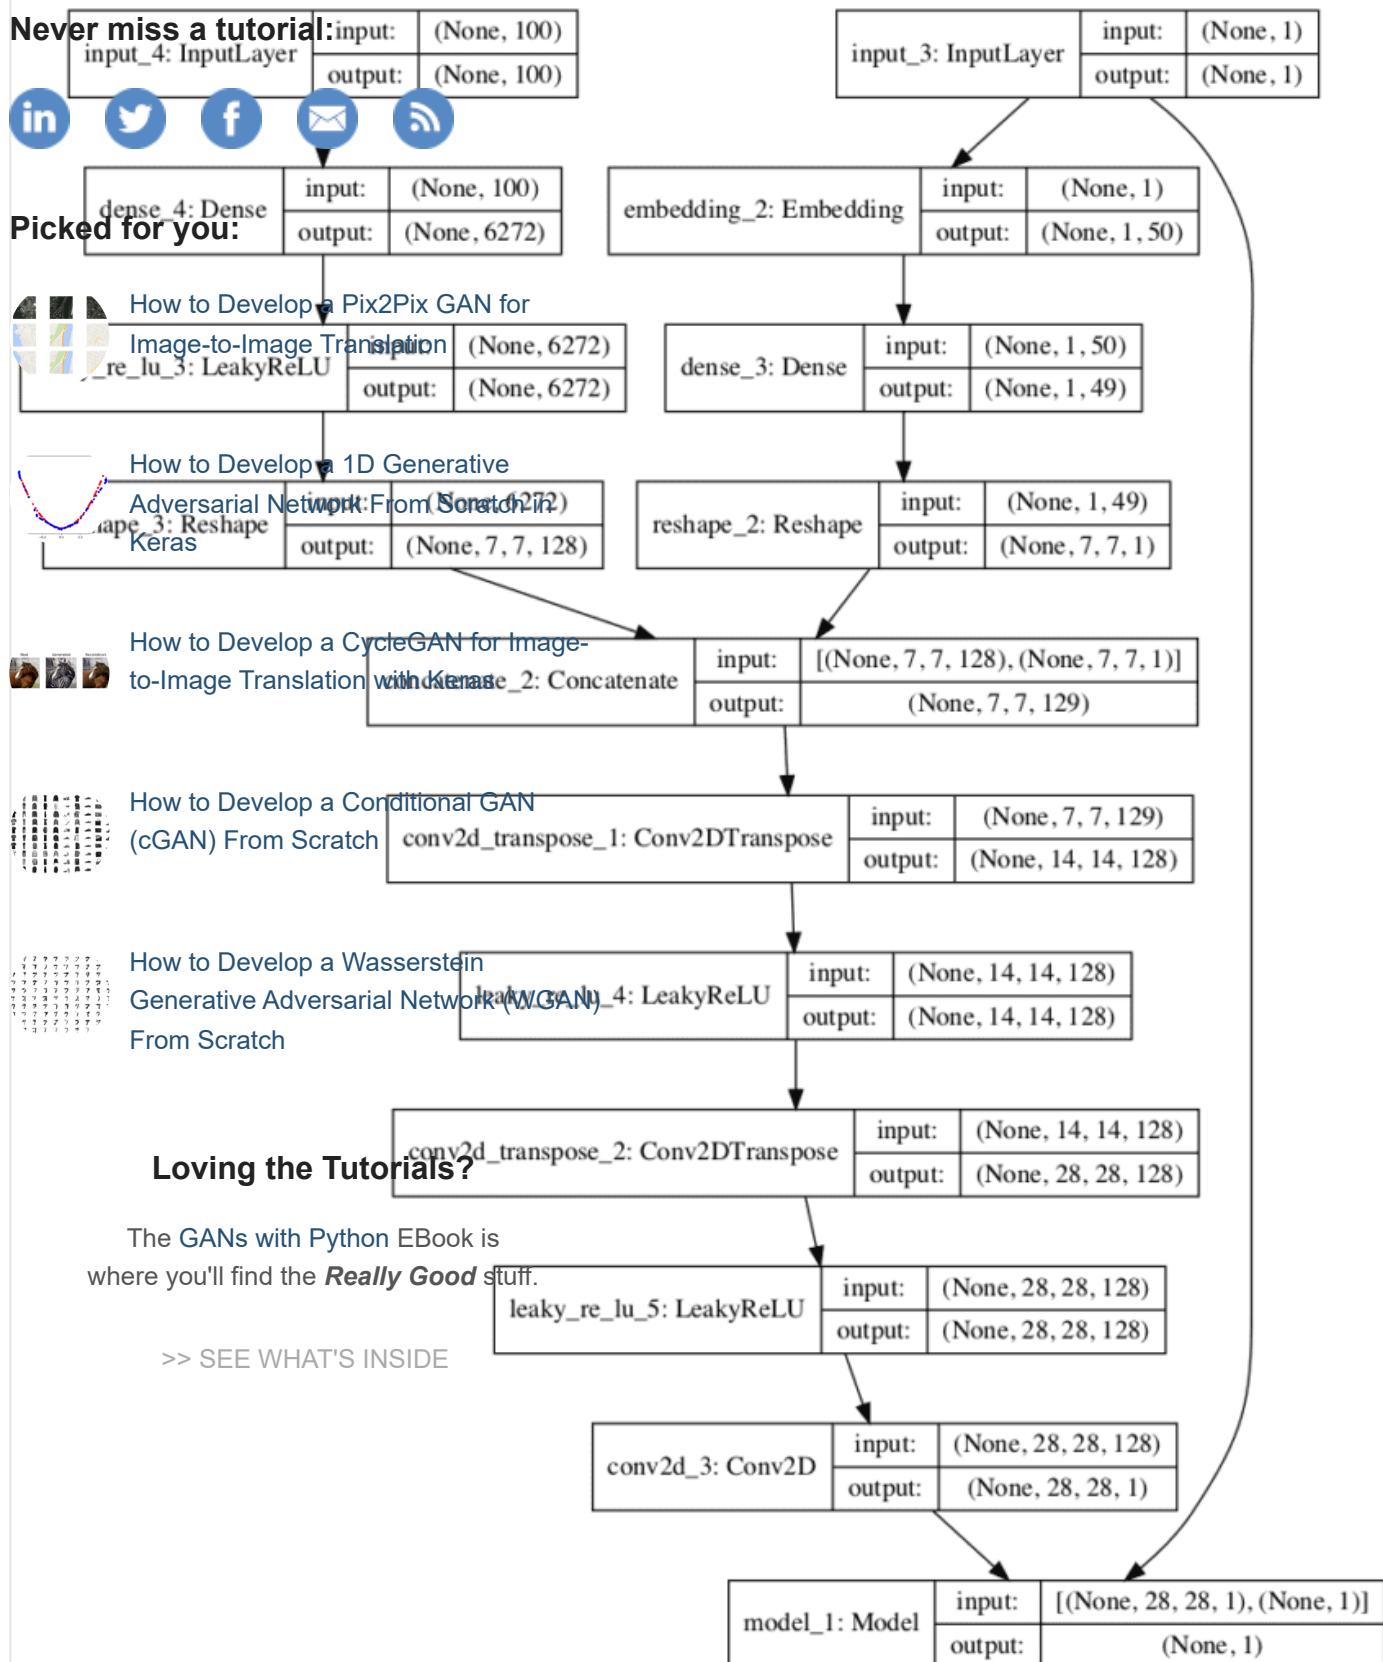
How to Develop a Conditional GAN  
The code below summarizes the composite GAN model.  
(cGAN) From Scratch

Importantly, it shows the generator model in full with the point in latent space and class label as input, and the connection of the output of the generator and the same class label as input to the discriminator model. The discriminator model is defined as taking the image output from the generator and the class label as input, and outputting a single class label classification of real or fake. From Scratch

## Loving the Tutorials?

The [GANs with Python](#) EBook is  
where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE



Plot of the Composite Generator and Discriminator Model in the Conditional Generative Adversarial Network

The hard part of the conversion from unconditional to conditional GAN is done, namely the definition and configuration of the model architecture.

Next, all that remains is to update the training process to also use class labels.

First, the `load_real_samples()` and `generate_real_samples()` functions for loading the dataset and selecting a batch of samples respectively must be updated to make use of the real class labels from the

training dataset. Importantly, the `generate_real_samples()` function now returns images, clothing labels, and the class label for the discriminator (class=1).

```

1 # load fashion mnist images
2 def load_real_samples():
3     # load dataset
4     (trainX, trainy), _, _ = load_data()
5     # expand to 3d, e.g. add channels
6     X = expand_dims(trainX, axis=-1)
7     # convert from ints to floats
8     X = X.astype('float32')
9     # scale from [0,255] to [-1,1]
10    X = (X - 127.5) / 127.5
11    return [X, trainy]
12
13 #/select real samples
14 def generate_real_samples(dataset, n_samples):
15     # split into images and labels
16     images, labels = dataset
17     # choose random instances
18     ix = randint(0, images.shape[0], n_samples)
19     # select images and labels
20     X, labels = images[ix], labels[ix]
21     # generate class labels
22     y = ones((n_samples, 1))
23     return [X, labels]

```

The `generate_latent_points()` function must be updated to also generate an array of randomly selected integer class labels to go along with the randomly selected points in the latent space.

The `generate_fake_samples()` function must be updated to use these randomly generated class labels as input to the generator model when generating new fake images.

```

1 # generate points in latent space as input for the generator
2 def generate_latent_points(latent_dim, n_samples, n_classes=10):
3     # generate points in the latent space
4     x_input = randn(latent_dim * n_samples)
5     # reshape into a batch of inputs for the network
6     z_input = x_input.reshape(n_samples, latent_dim)
7     # generate labels
8     labels = randint(0, n_classes, n_samples)
9     return [z_input, labels]
10
11 # use the generator to generate n fake examples, with class labels
12 def generate_fake_samples(generator, latent_dim, n_samples):
13     # generate points in latent space
14     z_input, labels_input = generate_latent_points(latent_dim, n_samples)
15     # predict outputs
16     images = generator.predict([z_input, labels_input])
17     # create class labels
18     y = zeros((n_samples, 1))
19     return [images, labels_input], y

```

Finally, the `train()` function must be updated to retrieve and use the class labels in the calls to updating the discriminator and generator models.

```

1 # train the generator and discriminator
2 def train(g_model, d_model, gan_model, dataset, latent_dim, n_epochs=100, n_batch=128):
3     bat_per_ago = int(dataset[0].shape[0] / n_batch)
4     half_batch = int(n_batch / 2)
5     # manually enumerate epochs
6     for i in range(n_epochs):
7         # enumerate batches over the training set
8         for j in range(bat_per_ago):
9             # get randomly selected 'real' samples

```

```

10 [X_real, labels_real], y_real = generate_real_samples(dataset, half_batch)
11 # update discriminator model weights
12 d_loss1, _ = d_model.train_on_batch([X_real, labels_real], y_real)
13 # generate 'fake' examples
14 [X_fake, labels], y_fake = generate_fake_samples(g_model, latent_dim, half_batch)
15 # update discriminator model weights
16 d_loss2, _ = d_model.train_on_batch([X_fake, labels], y_fake)
17 # prepare points in latent space as input for the generator
18 [z_input, labels_input] = generate_latent_points(latent_dim, n_batch)
19 # create inverted labels for the fake samples
20 y_gan = ones((n_batch, 1))
21 # update the generator via the discriminator's error
22 g_loss = gan_model.train_on_batch([z_input, labels_input], y_gan)
23 # summarize loss on this batch
24 print('>%d, %d/%d, d1=%f, d2=%f, g=%f' %
25 (i+1, half_batch_per_epoch, d_loss1, d_loss2, g_loss))
26 # save the generator model
27 g_model.save('cgan_generator.h5')

```

Keras

Tying all of this together, the complete example of a conditional deep convolutional generative adversarial network for the Fashion MNIST dataset is listed below.

[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

```

1 # example of training an conditional gan on the fashion mnist dataset
2 from numpy import expand_dims
3 from numpy import zeros
4 from numpy import ones
5 from numpy.random import randn
6 from numpy.random import randint
7 from keras.datasets.fashion_mnist import load_data
8 from keras.optimizers import Adam
9 from keras.models import Model
10 from keras.layers import Input
11 from GenerativeAdversarialNetworks(WGAN)
12 from keras.layers import Reshape
13 from keras.layers import Flatten
14 from keras.layers import Conv2D
15 from keras.layers import Conv2DTranspose
16 from keras.layers import LeakyReLU
17 from keras.layers import Dropout
18 from keras.layers import Embedding
19 from keras.layers import Concatenate
20
21 # define the standalone discriminator model
22 def define_discriminator(in_shape=(28,28,1), n_classes=10):
23     # label input
24     in_label = Input(shape=(1,))
25     # embedding for categorical input
26     li = Embedding(n_classes, 50)(in_label)
27     # scale up to image dimensions with linear activation
28     n_nodes = in_shape[0] * in_shape[1]
29     li = Dense(n_nodes)(li)
30     # reshape to additional channel
31     li = Reshape((in_shape[0], in_shape[1], 1))(li)
32     # image input
33     in_image = Input(shape=in_shape)
34     # concat label as a channel
35     merge = Concatenate()([in_image, li])
36     # downsample
37     fe = Conv2D(128, (3,3), strides=(2,2), padding='same')(merge)
38     fe = LeakyReLU(alpha=0.2)(fe)
39     # downsample
40     fe = Conv2D(128, (3,3), strides=(2,2), padding='same')(fe)
41     fe = LeakyReLU(alpha=0.2)(fe)
42     # flatten feature maps
43     fe = Flatten()(fe)
44     # dropout
45     fe = Dropout(0.4)(fe)
46     # output

```

```

47 out_layer = Dense(1, activation='sigmoid')(fe)
48 # define model
49 model = Model([in_image, in_label], out_layer)
50 # compile model
51 opt = Adam(lr=0.0002, beta_1=0.5)
52 model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
53 return model
54
55 # define the standalone generator model
56 def define_generator(latent_dim, n_classes=10):
57     # Label Input
58     in_label = Input(shape=(1,))
59     # embedding for categorical input
60     li = Embedding(n_classes, 50)(in_label)
61     # linear multiplication
62     n_nodes = 7 * 7
63     li = Dense(n_nodes)(li)
64     # reshape to additional channel
65     li = Reshape((7, 7, 1))(li)
66     # image generator input
67     in_lat = Input(shape=(latent_dim,))
68     # foundation for 7x7 image
69     n_nodes = 128 * 7 * 7
70     gen = Dense(n_nodes)(in_lat)
71     gen = LeakyReLU(alpha=0.2)(gen)
72     gen = Reshape((7, 7, 128))(gen)
73     # merge image gen and label input
74     merge = Concatenate()([gen, li])
75     # upsample to 14x14
76     gen = Conv2DTranspose(128, (4,4), strides=(2,2), padding='same')(merge)
77     gen = LeakyReLU(alpha=0.2)(gen)
78     # upsample to 28x28
79     gen = Conv2DTranspose(128, (4,4), strides=(2,2), padding='same')(gen)
80     gen = LeakyReLU(alpha=0.2)(gen)
81     # output
82     out_layer = Conv2D(1, (7,7), activation='tanh', padding='same')(gen)
83     # define model
84     model = Model([in_lat, in_label], out_layer)
85     return model
86
87 # define the combined generator and discriminator model, for updating the generator
88 def define_gan(g_model, d_model):
89     # make weights in the discriminator not trainable
90     d_model.trainable = False
91     # get noise and label inputs from generator model
92     gen_noise, gen_label = g_model.input
93     # get image output from the generator model
94     gen_output = g_model.output
95     # connect image output and label input from generator as inputs to discriminator
96     gan_output = d_model([gen_output, gen_label])
97     # define gan model as taking noise and label and outputting a classification
98     model = Model([gen_noise, gen_label], gan_output)
99     # compile model
100    opt = Adam(lr=0.0002, beta_1=0.5)
101    model.compile(loss='binary_crossentropy', optimizer=opt)
102    return model
103
104 # load fashion mnist images
105 def load_real_samples():
106     # load dataset
107     (trainX, trainy), (_, _) = load_data()
108     # expand to 3d, e.g. add channels
109     X = expand_dims(trainX, axis=-1)
110     # convert from ints to floats
111     X = X.astype('float32')
112     # scale from [0,255] to [-1,1]
113     X = (X - 127.5) / 127.5
114     return [X, trainy]
115

```

```

116 # # select real samples
117 # miss a tutorial: def generate_real_samples(dataset, n_samples):
118     # split into images and labels
119     images, labels = dataset
120     # choose random instances
121     ix = randint(0, images.shape[0], n_samples)
122     # select images and labels
123     X, labels = images[ix], labels[ix]
124     # generate class labels
125     y = ones((n_samples, 1))
126     return [X, labels], y
127 
128     # Image-to-Image Translation
129     # generate points in latent space as input for the generator
130     def generate_latent_points(latent_dim, n_samples, n_classes=10):
131         # generate points in the latent space
132         x_input = randn(latent_dim * n_samples)
133         # reshape into a batch of inputs for the network
134         # Adversarial network from scratch
135         z_input = x_input.reshape(n_samples, latent_dim)
136         # generate labels
137         labels = randint(0, n_classes, n_samples)
138         return [z_input, labels]
139 
140         How to Develop a CycleGAN for Image-
141         # use the generator to generate n fake examples, with class labels
142         def generate_fake_samples(generator, latent_dim, n_samples):
143             # generate points in latent space
144             z_input, labels_input = generate_latent_points(latent_dim, n_samples)
145             # predict outputs
146             images = generator.predict([z_input, labels_input])
147             # create class labels
148             y = zeros((n_samples, 1))
149             return [images, labels_input], y
150 
151         # train the generator and discriminator
152         def train(gan_model, d_model, gan_model, dataset, latent_dim, n_epochs=100, n_batch=128):
153             bat_per_epo = int(dataset[0].shape[0] / n_batch)
154             half_batch = int(n_batch / 2)
155             # manually enumerate epochs
156             for i in range(n_epochs):
157                 # enumerate batches over the training set
158                 for j in range(bat_per_epo):
159                     # get randomly selected 'real' samples
160                     [X_real, labels_real], y_real = generate_real_samples(dataset, half_batch)
161                     # update discriminator model weights
162                     d_loss1, _ = d_model.train_on_batch([X_real, labels_real], y_real)
163                     # generate 'fake' examples
164                     [X_fake, labels_fake], y_fake = generate_fake_samples(g_model, latent_dim, half_batch)
165                     # update discriminator model weights
166                     d_loss2, _ = d_model.train_on_batch([X_fake, labels_fake], y_fake)
167                     # prepare points in latent space as input for the generator
168                     [z_input, labels_input] = generate_latent_points(latent_dim, n_batch)
169                     # create inverted labels for the fake samples
170                     y_gan = ones((n_batch, 1))
171                     # update the generator via the discriminator's error
172                     g_loss = gan_model.train_on_batch([z_input, labels_input], y_gan)
173                     # summarize loss on this batch
174                     print('>%d, %d/%d, d1=%f, d2=%f, g=%f' %
175                         (i+1, j+1, bat_per_epo, d_loss1, d_loss2, g_loss))
176                     # save the generator model
177                     g_model.save('cgan_generator.h5')
178 
179                     # size of the latent space
180                     latent_dim = 100
181                     # create the discriminator
182                     d_model = define_discriminator()
183                     # create the generator
184                     g_model = define_generator(latent_dim)
185                     # create the gan
186                     gan_model = define_gan(g_model, d_model)
187                     # load image data

```

```

185 dataset = load_real_samples()
186 # train model
187 train(g_model, d_model, gan_model, dataset, latent_dim)

```

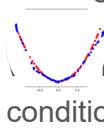
LinkedIn  Twitter  Example  takes some time, and GPU hardware is recommended, but not required.

At the end of the run, the model is saved to the file with name ‘*cgan\_generator.h5*’.

## Conditional Clothing Generation

 Image-to-Image Translation

In this section, we will use the trained generator model to conditionally generate new photos of items of clothing.

 How to Develop a 1D Generative Adversarial Network From Scratch in Keras

Update our code example for generating new images with the model to now generate images conditional on the class label. We can generate 10 examples for each class label in columns.

The complete example is listed below:

 Image-to-Image Translation with Keras

```

1 # example of loading the generator model and generating images
2 from numpy import asarray
3 from numpy.random import randn
4 from numpy.random import randint
5 from keras.models import load_model
6 from matplotlib import pyplot
7
8 # generate points in latent space as input for the generator
9 def generate_latent_points(latent_dim, n_samples, n_classes=10):
10    # generate points in the latent space
11    x_input = randn(latent_dim * n_samples)
12    # reshape into a batch of inputs for the network
13    z_input = x_input.reshape(n_samples, latent_dim)
14    # generate labels
15    labels = randint(0, n_classes, n_samples)
16    return [z_input, labels]
17
18 # create and save a plot of generated images
19 def save_plot(examples, n):
20    # plot images
21    for i in range(n * n):
22        # define subplot
23        pyplot.subplot(n, n, 1 + i)
24        # turn off axis
25        pyplot.axis('off')
26        # plot raw pixel data
27        pyplot.imshow(examples[i, :, :, 0], cmap='gray_r')
28    pyplot.show()
29
30 # load model
31 model = load_model('cgan_generator.h5')
32 # generate images
33 latent_points, labels = generate_latent_points(100, 100)
34 # specify labels
35 labels = asarray([x for _ in range(10) for x in range(10)])
36 # generate images
37 X = model.predict([latent_points, labels])
38 # scale from [-1,1] to [0,1]
39 X = (X + 1) / 2.0
40 # plot the result
41 save_plot(X, 10)

```

Running the example loads the saved conditional GAN model and uses it to generate 100 items of clothing.

The clothing is organized in columns. From left to right, they are “*t-shirt*”, ‘trouser’, ‘pullover’, ‘dress’, ‘coat’, ‘sandal’, ‘shirt’, ‘sneaker’, ‘bag’, and ‘ankle boot’.

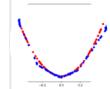


You can see not only all the randomly generated items of clothing plausible, but they also match their expected category.

### Picked for you:



[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)



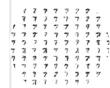
[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)



[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)



[How to Develop a Conditional GAN \(cGAN\) From Scratch](#)



[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)



**Loving the Tutorials?**

The [GANs with Python EBook](#) is where you'll find the **Really Good** stuff.

Example of 100 Generated items of Clothing using a Conditional GAN.

[>> SEE WHAT'S INSIDE](#)

## Extensions

This section lists some ideas for extending the tutorial that you may wish to explore.

- **Latent Space Size.** Experiment by varying the size of the latent space and review the impact on the quality of generated images.
- **Embedding Size.** Experiment by varying the size of the class label embedding, making it smaller or larger, and review the impact on the quality of generated images.
- **Alternate Architecture.** Update the model architecture to concatenate the class label elsewhere in the generator and/or discriminator model, perhaps with different dimensionality, and review the impact on the quality of generated images.

If you explore any of these extensions, I'd love to know.

Post your findings in the comments below.

## Further Reading



see more resources on the topic if you are looking to go deeper.

### Posts

#### Picked for you:

- Chapter 20. Deep Generative Models, Deep Learning, 2016.



### Papers



Adversarial Network From Scratch in Keras.

- Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, 2015



Image-To-Image Translation With Conditional Adversarial Networks, 2017.

- Conditional Generative Adversarial Nets For Convolutional Face Generation, 2015.



- Keras Datasets API.
- Keras Sequential Model API
  - How to Develop a Wasserstein GAN
- Keras Convolutional Layers API
  - Generative Adversarial Network (WGAN)
  - How can I “freeze” Keras layers?
- Matplotlib API
- NumPy Random sampling (numpy.random) API
- NumPy Array manipulation routines

### Loving the Tutorials?

### Articles

The GANs with Python EBook is where you'll find the **Really Good** stuff.

- How to Train a GAN? Tips and tricks to make GANs work
- Fash >> SEE WHAT'S INSIDE b.
- Training a Conditional DC-GAN on CIFAR-10 (code), 2018.
- GAN: From Zero to Hero Part 2 Conditional Generation by GAN, 2018.
- Keras-GAN Project. Keras implementations of Generative Adversarial Networks, GitHub.
- Conditional Deep Convolutional GAN (CDCGAN) – Keras Implementation, GitHub.

## Summary

In this tutorial, you discovered how to develop a conditional generative adversarial network for the targeted generation of items of clothing.

Specifically, you learned:

- The limitations of generating random samples with a GAN that can be overcome with a conditional generative adversarial network.
- How to develop and evaluate an unconditional generative adversarial network for generating photos of items of clothing.

• How to develop and evaluate a conditional generative adversarial network for generating photos of items of clothing.



Ask your questions in the comments below and I will do my best to answer.

**Picked for you:**

How to Develop a Pix2Pix GAN for Image-to-Image Translation

## Develop Generative Adversarial Networks Today!



### Develop Your GAN Models in Minutes

...with just a few lines of python code

Discover how in my new Ebook:  
Generative Adversarial Networks with Python

It provides **self-study tutorials** and **end-to-end projects** on:  
*DCGAN, conditional GANs, image translation, Pix2Pix, CycleGAN*  
and much more...

### Finally Bring GAN Models to your Vision Projects

Skip the Academics. Just Results.

[SEE WHAT'S INSIDE](#)

### Loving the Tutorials?

[Tweet](#)

[Tweet](#)

[Share](#)

[Share](#)

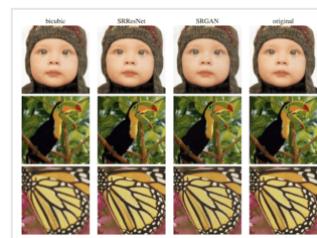
The [GANs with Python](#) EBook is where you find the *Really Good* stuff.

### More On This Topic

[>> SEE WHAT'S INSIDE](#)



How to Develop an Auxiliary Classifier GAN (AC-GAN)...



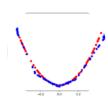
18 Impressive Applications of Generative Adversarial...

**Never miss a tutorial:****Picked for you:**

A Gentle Introduction to Pix2Pix Generative...



[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)



[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

A Tour of Generative Adversarial Network Models



[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

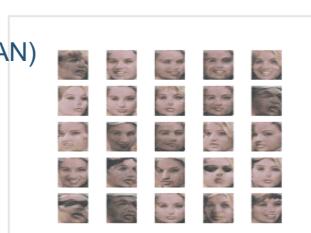


[How to Develop a Conditional GAN \(cGAN\) From Scratch](#)

How to Implement a Semi-Supervised GAN (SGAN) From...



[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)

**Loving the Tutorials?**The [GANs with Python](#) Ebookswhere you'll find the **Really Good** stuff.>> SEE WHAT'S INSIDE **ownlee**

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

&lt; How to Explore the GAN Latent Space When Generating Faces

How to Identify and Diagnose GAN Failure Modes &gt;

183 Responses to *How to Develop a Conditional GAN (cGAN) From Scratch*

Keith Beaudoin July 5, 2019 at 8:52 am #

REPLY ↗

cool ty  
**Never miss a tutorial:**



 **Jason Brownlee** July 5, 2019 at 8:54 am #

REPLY ↗

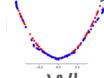
**Picked for you:**

Thanks Keith.



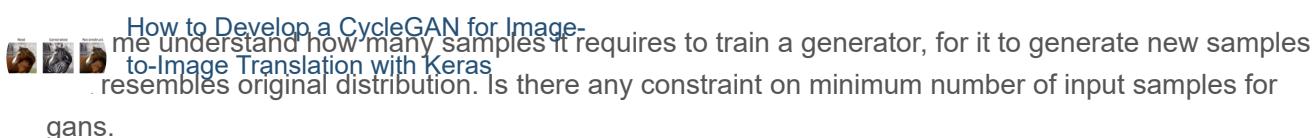
 **Shravan Kumar Parunandula** July 5, 2019 at 11:47 am #

REPLY ↗

 **How to Develop a 1D Generative Adversarial Network From Scratch in Keras**

This is fantastic. Thanks for disseminating great knowledge.

What if I wanted to train the discriminator as well, as we are only training generator in the current model. Please correct me if I am wrong.



 **How to Develop a CycleGAN for Image-to-Image Translation with Keras**

me understand how many samples it requires to train a generator, for it to generate new samples resembles original distribution. Is there any constraint on minimum number of input samples for gans.



 **How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch**

REPLY ↗

No, both the generator and discriminator are trained at the same time.

There is great work with the semi-supervised GAN on training a classifier with very few real samples.  
**Loving the Tutorials?**

The [GANs with Python](#) EBook is where you'll find the **Really Good** stuff.

 **Shabnam** July 7, 2019 at 9:40 pm #

REPLY ↗

>> SEE WHAT'S INSIDE

NICE BLOG.

Do you have any blog on deployment of pytorch or tensorflow based gan model on Android? I am desperately in need of it.



**Jason Brownlee** July 8, 2019 at 8:41 am #

REPLY ↗

No, sorry.



**Raja** August 9, 2020 at 1:04 am #

REPLY ↗

Hi Jason , But you are setting the discriminator weights trainable as False  
# make weights in the discriminator not trainable  
d\_model.trainable = False

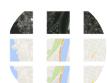
I don't understand it now .  
**Never miss a tutorial:**



**Jason Brownlee** August 9, 2020 at 5:45 am #

REPLY ↗

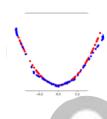
**Picked for you:** Only in the context of the composite model.



To learn more about freezing weights in different contexts, see:

[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#) – How can I freeze layers and do fine-tuning?

[https://keras.io/getting\\_started/faq/](https://keras.io/getting_started/faq/)

 How to Develop a 1D Generative Adversarial Network From Scratch in Keras  
**Ken** July 5, 2019 at 1:01 pm #

REPLY ↗

 A ton of great blog posts! I'm really excited for your book on GAN's. I think bugged you about image-to-image translation with Keras.  
**Partha S** July 5, 2019 at 4:23 pm #

 How to Develop a Conditional GAN  
**Jason Brownlee** July 6, 2019 at 8:21 am #

REPLY ↗

Thanks! And thanks for bugging me Ken!

 I'm really excited about it.  
 How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch  
**Ken**

REPLY ↗

 **Partha S** July 5, 2019 at 4:23 pm #  
**Loving the Tutorials?**

Ken  
 The [GANs with Python](#) EBook is  
 where you'll find the [Really Good](#) stuff here please

Regards  
**Partha** >> SEE WHAT'S INSIDE



**Jason Brownlee** July 6, 2019 at 8:24 am #

REPLY ↗

Ken is referring to my upcoming book on GANs.

The title will be "Generative Adversarial Networks with Python".

It should be available in a week or two.



**Hitarth** July 5, 2019 at 6:24 pm #

REPLY ↗

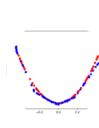
I didn't understand that how the generator will produce good results while training composite  
**Never miss a tutorial:**  
 unconscious GAN by passing ones as output label, shouldn't it be zeros?



**Picked for you!** **Hitarth** July 5, 2019 at 6:25 pm #

REPLY ↗

 In the unconditional GAN training.  
[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)

 **Jason Brownlee** July 6, 2019 at 8:31 am #  
 How to Develop an Adversarial Network From Scratch in Keras  
 The unconditional GAN is trained.

REPLY ↗

Perhaps I don't understand your question?  
[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

 **Jason Brownlee** July 6, 2019 at 8:31 am #  
 How to Develop a Conditional GAN (cGAN) From Scratch

REPLY ↗

Basically, we are training the generator to fool the discriminator, and in this case, the generator is conditional on the specific class label. The discriminator causes the discriminator to associate specific generated images with class labels.  
[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\)](#)

If [From Scratch](#) to you, perhaps start here:

<https://machinelearningmastery.com/what-are-generative-adversarial-networks-gans/>

## Loving the Tutorials?

 **Yasser** March 22, 2020 at 2:14 am #  
 where you'll find the **Really Good** stuff.

REPLY ↗

Hello sir  
 I re >> SEE WHAT'S INSIDE have a question, I want to work with this technique but as an input I have an Image and then I feed it to the generator to have another image and then feed it to the discriminator but the problem is all tutorials are starting from a random input .  
 Do you have any blog or code you can help me with

 **Jason Brownlee** March 22, 2020 at 6:58 am #

REPLY ↗

It sounds like you might be interested in image to image translation.

This will help:

<https://machinelearningmastery.com/a-gentle-introduction-to-pix2pix-generative-adversarial-network/>

**Howard** July 12, 2019 at 9:31 am #

REPLY ↗

Great article, thank you! I have two questions.

### Never miss a tutorial:

First, you use an embedding layer on the labels in both the discriminator and generator. I don't see the need, the embedding is useful for you. With just 10 labels, why is a 50-dimensional vector any more useful than a normal one-hot vector (after all, the ten one-hot vectors are orthonormal, so they're as distinct as can be). So what is the algorithmic motivation for having an embedding layer?

### Picked for you:

Second, why then follow that with a dense layer? Again, the one-hot label vectors seem to be all we need, [How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#) already does image Translation

Thank you!

[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

 Jason Brownlee July 13, 2019 at 6:47 am #

REPLY ↗

[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

The embedding layer provides a projection of the class label, a distributed representation that can be used to condition the image generation and classification.

[How to Develop a Conditional GAN \(cGAN\) From Scratch](#)

There are other ways of getting the class label into the model, but this approach is reported to be more effective. Why? This is a hard question and might be intractable right now. Most of the GANs in [Generative Adversarial Network \(WGAN\) From Scratch](#)

Try the alternate of just a one hot encoded vector concat with the z for G and a secondary input for D and compare results.

### Loving the Tutorials?

The [GANs with Python EBook](#) is where you'll find the [More Really Good stuff](#).

REPLY ↗

>> SEE WHAT'S INSIDE

Thank you for this very useful and detailed. Do you have any references that explain the embedding idea more thoroughly, or can you offer any more intuition? I understand why you might use an embedding for words/sentences as there is an idea of semantic similarity there, but not following why in a dataset like this (or simple MNIST) an embedding layer makes sense. Is it effectively just a way of reshaping the one-hot? Thanks!



 Jason Brownlee October 23, 2019 at 1:49 pm #

REPLY ↗

An embedding is an alternate to one hot encoding for categorical data.

It is popular for words, but can be used for any categorical or ordinal data.

 Chuanliang Jiang July 14, 2019 at 7:56 am #

REPLY ↗

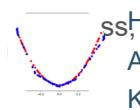
**Never miss a tutorial:** In unconditional GAN codes, why discriminator model weights can be updated separately for exclusive real and fake sample ?

up dis nat odesights  
 d\_loss1, \_ = d\_model.train\_on\_batch(X\_real, y\_real)

# update discriminator model weights

d\_loss2, \_ = d\_model.train\_on\_batch(X\_fake, y\_fake)

 How to Develop a Pix2Pix GAN for  
 Specifically discriminator is binary classification. If all samples are real (=1) or faked(=0) exclusively , the  
 Image-to-Image Translation  
 binary classification is unable to be converged. Why not combined X\_real and X\_fake together and then  
 input the sample into discriminator which will classify real and faked sample, e.g.

 ss, \_ = d\_model.train\_on\_batch([X\_real, X\_fake], [y\_real, y\_fake])  
 Adversarial Network From Scratch in  
 Keras

 Jason Brownlee August 14, 2019 at 8:19 am #  
 REPLY ↗

How to Develop a CycleGAN for Image

Translation with Keras

You can, but it has been reported that separate batch updates keep the D model stable with respect to the performance of the generator (e.g. it does not get better – faster).

 How to Develop a Conditional GAN  
 (cGAN) From Scratch  
<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

 Kristof August 30, 2019 at 12:02 am #  
 REPLY ↗

How to Develop a Wasserstein  
 Generative Adversarial Network (WGAN)

From Scratch

Thanks for the very useful tutorial!

I always get these kind of warnings:

**Loving the Tutorials?**

"W0829 11:18:47.925395 14568 training.py:2197] Discrepancy between trainable weights and collected trainable weights, did you set model.trainable without calling model.compile after ?"

where you'll find the **Really Good** stuff.

Does it mean I did something different, or is this something you see as well. The model runs, so does it matter' >> SEE WHAT'S INSIDE

 Jason Brownlee August 30, 2019 at 6:25 am #  
 REPLY ↗

You can safely ignore that warning – we are abusing Keras a little 😊

 Yue September 8, 2019 at 10:42 pm #  
 REPLY ↗

Hi Janson, very nice tutorial< I was stuck somewhere when running your code:

we define "define\_discriminator(in\_shape=(28,28,1))" with shape (28,28,1), and then we call it to do "d\_model.train\_on\_batch(X\_real, y\_real)", where the sample size is 64 (error message as follows):

ValueError: Error when checking model input: the list of Numpy arrays that you are passing to your model is not the size the model expected. Expected to see 1 array(s), but instead got the following list of

64 arrays: [ ,  
**Never miss a tutorial:**

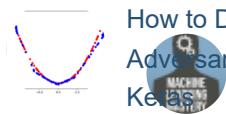
I am new to deep learning so do not know how to fix it..



**Picked for you:** **Yue** September 9, 2019 at 2:12 am #

REPLY ↗

 Sorry, I found out I made a mistake when I tried to copy your code< Ignore my question  
 How to Develop a Pix2Pix GAN for  
 Image-to-Image Translation

 How to Develop a 1D Generative  
 Adversarial Keras **Jason Brownlee** September 9, 2019 at 5:16 am #

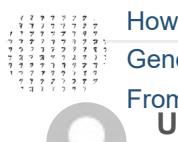
REPLY ↗

No problem!

 How to Develop a CycleGAN for Image-  
 to-Image Translation with Keras

REPLY ↗

 How to Develop a Conditional GAN  
 (cGAN) From Scratch Sorry to hear that, I have some suggestions here that might help:  
<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me>

 How to Develop a Wasserstein  
 Generative Adversarial Network (WGAN)  
 From Scratch **Umair Khan** September 26, 2019 at 10:01 pm #

REPLY ↗

How is the discriminator model instance 'd\_model' trained in the training loop, when the same  
**Loving the Tutorials?** instance is set to trainable=False in the 'define\_GAN' method?

The **GANs with Python** EBook is  
 where you'll find the **Really Good** stuff.

 >> SEE WHAT'S INSIDE September 27, 2019 at 8:01 am #

REPLY ↗

Setting trainable=False only effects the generator, it does not effect the discriminator.

See this:

<https://keras.io/getting-started/faq/#how-can-i-freeze-keras-layers>

 **W. Jin** October 19, 2019 at 7:23 am #

REPLY ↗

Thank you very much! This is a clearly written, nicely structured and most inspiring tutorial. I  
 love all the detailed explanations as well as the attached code. Excellent!

 **Jason Brownlee** October 20, 2019 at 6:12 am #

REPLY ↗

Thanks!  
Never miss a tutorial:



Venugopal Shah October 22, 2019 at 3:09 pm #

REPLY ↗

### Picked for you:

Thank you Jason! This is a really good tutorial. I went on to do some further experiments and I'm facing some issues. I am trying to implement a cSAGAN with spectral normalization and for some reason, the discriminator throws up an error 'NoneType' object has no attribute '\_inbound\_nodes'.

[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)  
[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

I need some insight from you on what could be wrong.

[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

Jason Brownlee October 23, 2019 at 6:30 am #

REPLY ↗

Not sure I can help you with your custom code, sorry. Perhaps try posting to stackoverflow?  
[How to Develop a Conditional GAN \(cGAN\) From Scratch](#)

[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\)](#)

I managed to figure it out eventually! It was a really silly mistake on my part. However I would still like to thank you again for this post which was the founding base for my project. You are doing a wonderful job!

### Loving the Tutorials?

The [GANs with Python](#) EBook is

where you'll find the   
 Jason Brownlee October 27, 2019 at 5:37 am #

REPLY ↗

>> SEE WHAT'S INSIDE r that, well done!



Jordan November 10, 2019 at 6:58 am #

REPLY ↗

Hi,

Does using an Embedding layer segment the latent space for each class? What I mean is can you use this method to get the generator to produce, for instance, a t-shirt+shoe combination?



Jason Brownlee November 10, 2019 at 8:28 am #

REPLY ↗

Yes! Probably.

**Never miss a tutorial:**

Dang Tuan Hoang December 3, 2019 at 1:14 pm #

REPLY ↗



Do you think it is possible to train Conditional GAN with more than 1 condition? For example, I want to **Picked for you** which color the input picture, conditioned by non-color input picture and color label



How to Develop a Pix2Pix GAN for

Image-to-Image Translation



Jason Brownlee December 3, 2019 at 1:36 pm #

REPLY ↗

How to Develop a 1D Generative

Adversarial Network From Scratch in

Yes Keras

multiple conditions.

I think Infogan has

Yes sounds like a little image to image translation and a little conditional gan. Give it a try! Let me

know how you go.



How to Develop a CycleGAN for Image-

to-Image Translation with Keras



Dang Tuan Hoang December 3, 2019 at 3:35 pm #

REPLY ↗



How to Develop a Conditional GAN

(cGAN) From Scratch

I'm currently following MC-GAN paper, but Info GAN look interesting as well. Definitely,

gonna try it later



How to Develop a Wasserstein

Generative Adversarial Network (WGAN)

From Scratch



Jason Brownlee December 4, 2019 at 5:28 am #

REPLY ↗

Very cool. Eager to hear how you go.

**Loving the Tutorials?**

The GANs with Python EBook is

Marcin December 17, 2019 at 6:40 pm #

where you'll find the **Really Good** stuff.

REPLY ↗

! Thank you for your effort.

&gt;&gt; SEE WHAT'S INSIDE



Jason Brownlee December 18, 2019 at 6:00 am #

REPLY ↗

You're welcome.



tvtaerum January 4, 2020 at 9:36 am #

REPLY ↗

As always, some beautiful work Brownlee. I realize this will be no surprise to you but you can run an almost identical cgan against the mnist (number) dataset by changing the following lines:

1. replace:

from tensorflow.keras.datasets.fashion\_mnist import load\_data

with:

from tensorflow.keras.datasets.mnist import load\_data

Note: I personally use the tensorflow version of keras

## Never miss a tutorial:

2. replace all instances of:

`opt = Adam(lr=0.0002, beta_1=0.5)`  
with:  


`opt = Adam(lr=0.0001, beta_1=0.5) # learning rate just needs to be slowed down`

## Picked for you:

The reason I mention this (almost obvious) fact is for me, cgan produces better and more interesting results than a simple gan. Other people's tutorials give some code and then with a wave of the hands How to Develop a Pix2Pix GAN for Image-to-Image Translation test that, "it's apparent something recognizable as numbers are being produced".

In particular I appreciated your explanations for the Keras functional API approach.

What I do wish is there was some easy way to graphically illustrate the equivalent of first and second derivative estimates in order to better "see" why some attempts fail and other succeed. I realize that an approximation to an approximation is difficult to visualize but something along those lines would be great since single number measures tell me almost nothing diagnostic about what's going on internally. For instance, (for illustration purposes only), it might be doing well with noses and chins but doing a poor job How to Develop a CycleGAN for Image-to-Image Translation with Keras eyes and ears. I'm sure there are many searching for better ways to have more science and less art these weight estimations. I'm sure you have some great insights on this.

Again, beautiful work and thank you for your great explanations.

How to Develop a Conditional GAN (cGAN) From Scratch



**Jason Brownlee** January 5, 2020 at 7:01 am #

REPLY ↗

How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch

Thanks for your feedback and kind words.

Generative Adversarial Network (WGAN)

Evaluating GANs remains challenging, some of the metrics here might give you ideas:

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

## Loving the Tutorials?



The GANs with Python EBook is available on [tvtaerum](#) January 13, 2020 at 4:14 pm #

REPLY ↗

where you'll find the **Really Good** stuff.

Thanks for your reply and pointing me to your tutorial. You may consider me to be a matter >> SEE WHAT'S INSIDE >> see some of the things I've attempted but my interest is in "what works". If there are useful observations to make, I'm sure you'll do a much more elegant job than I can. My observations are based on a limited number of experiments – I am amazed at how much you accomplish every month.

Simply for interest sake, I have a set of learning rates and betas which consistently produce good results for both the MNIST and the FASHION\_MNIST dataset for me. I realize learning rates and betas are not the bleeding edge of GANS research but I am surprised the narrow range over which there is convergence:

`define_discriminator opt as:`

`opt = Adam(lr=0.0008, beta_1=0.1)`

`define_gan opt as:`

`opt = Adam(lr=0.0006, beta_1=0.05)`

I read your outline about measuring the goodness of results for gans. The tutorial is interesting but it seems to me you make your best point where you indicate that d1 and d2 ought to be about 0.6 while g ought to have values around 0.8. My general limited experience is, if I can keep my values for d1, d2 and

g within reasonable bounds of those values, then I am soon going to have convergence. And if I am **Never miss a tutorial:** going to have convergence, then I need to first obtain good estimates of learning rate and momentum.

 keep going with this, make point in your tutorial about exploration of latent space that you may have to restart the analysis if the values of loss go out of bounds. In keeping with this view, I attempted the following which I've added to the training function of your tutorial on exploring latent space.

**Picked for you:** Substantially, it saves a recent copy of the models where the values of loss are under 1.0 and recovers

the models when the losses go out of an arbitrary bound and "tries again". Surprisingly, it does seem to  How to Develop a Pix2Pix GAN for Image-to-Image Translation from where it left off and it does appear to prevent having to restart the whole analysis.

```

1  if (d_loss1 > 1.4 or d_loss2 > 1.4 or g_loss > 1.4):
2      qReSet = True
3      g_model = g_model_save
4      d_model = d_model_save
5      gan_model = gan_model_save
6      # summarize loss on this batch
7      if (j+1) % 5==0 or d_loss1 > 1.0 or d_loss2 > 1.0 or g_loss > 1.0:
8          diff = int(time.time()-now)
9          print('>%d/%d, %d/%d, d1=%f, d2=%f, g=%f, secs=%d' %
10              (i+1, n_epochs, j+1, bat_per_epo, d_loss1, d_loss2, g_loss, diff), qReSet)
11         if d_loss1 <= 1.0 and d_loss2 <= 1.0 and g_loss <= 1.0:
12             g_model_save = g_model
13             d_model_save = d_model
14             gan_model_save = gan_model

```

 understand what is the "maximum clarity" possible with respect to images generated. As in any statistical analysis, knowing the "maximum" is critical to understanding how far we've gotten or might theoretically go. While I recognize the mathematical usefulness of using normally distributed noise in the latent space, it doesn't appear to be important in practice – uniform noise between 3.0 and 3.0 (picky/kinky) works as well as normal. I've found the following works quite well:

From Scratch

```

1  initX = -3.0
2  rangeX = 2.0*abs(initX)
3  stepX = rangeX / (latent_dim * n_samples)
4  x_input = asarray([initX + stepX*(float(i)) for i in range(0,latent_dim * n_samples)])
5  shuffle(x_input)

```

The GANs with Python EBook is It's not terribly clever but it demonstrates, I think, that the points in the latent space do not have to be where you'll find the **Really Good** stuff random spaced but different and spread out, and there may be some benefit in insuring that the latent space is Gaussian. >> SEE WHAT'S INSIDE Gaussians themselves, whether or not this is the case in practice over a wide range of problems – you would obviously know better.

Finally, for the exploration of latent space, my GPU doesn't appear to have enough memory to use n\_batch = 128 so I'm using n\_batch = 64.

If I'm doing anything really dumb, feel free to let me know. 😊



**Jason Brownlee** January 14, 2020 at 7:19 am #

REPLY ↗

Very impressive, thanks for sharing!

You should consider writing up your findings more fully in a blog post.

**tvtaerum** January 16, 2020 at 4:47 am #

REPLY ↗



I apologize for putting so much in the comment section of your site. Your work is amazingly good and a person realizes this only after trying out different approaches and searching for better material on the Internet. My plan is, as you suggest, to put something up as a blog post to solve couple of issues.

But yes, I did do and report something really dumb in my last comment which I'd like to correct... I copied the address rather than using the 'copy' module and creating a backup. And, of course, I can only do this easily with the generator and the gan function. The compiled discriminator continues to work in background gradually improving on its ability to tell the difference between real and fake, while in the generator model jiggers its way to creating better fakes. In some ways this is how humans learn – the “slow” discriminator gradually improves over time, and the generator catches up.

Backing up the generator and gan models, I'm able to give every analysis many "second chances" Adversarial Network From Scratch in converging (not going out of bounds). In the interest of forcing convergence (irrespective of how really good the final model is) I used the following code:

Generative Adversarial Network (WGAN)  
put up a blog post and make many references to your great work once I better understand the  
limits.  
From Scratch

## Loving the Tutorials?



**Jason Brownlee** January 16, 2020 at 6:25 am #  
ANS with Python EBook is

REPLY ↵

where you'll find the ***Really Good*** stuff.  
Thanks for sharing.

>> SEE WHAT'S INSIDE



Sap February 10, 2020 at 11:43 pm #

REPLY ↵

GANs can be used for non-image data?



**Jason Brownlee** February 11, 2020 at 5:13 am #

REPLY ↵

Yes, but they are not as effective as specialized generative models, at least from what I have seen.



**marpuri\_ganesh** March 4, 2020 at 12:35 am #

REPLY ↶

your model fails while running it with mnist dataset but it works perfectly fine with fashion\_mnist dataset I

### Never miss a tutorial:

can not understand what is going wrong. Both d1\_loss and d2\_loss becomes 0.00 and gan\_loss

skyrockets could you give a hint in what's going wrong here



### Picked for you:



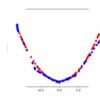
**Jason Brownlee** March 4, 2020 at 5:56 am #

REPLY ↗



How to Develop a Pix2Pix GAN for

If you change the dataset, you may need to tune the model to the change.



How to Develop a 1D Generative

Adversarial Network From Scratch

Keras

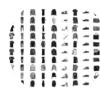
REPLY ↗

can you told me how to do that ? kindly

i build both model manually and bigger then this still both loss goes to zero

How to Develop a CycleGAN for Image-to-Image Translation with Keras

i don't know what i'm missing !!!!kindly help with this



How to Develop a Conditional GAN

(cGAN) From Scratch



**Jason Brownlee** July 20, 2020 at 6:14 am #

REPLY ↗



How to Develop a Generative Adversarial Network (GAN)

https://machinelearningmastery.com/start-here/#gans

From Scratch

### Loving the Tutorials?



September 9, 2021 at 7:10 am #

The GANs with Python EBook is Hello Jason, i got some problem with EMNIST dataset.

where you'll find the **Really Good** stuff.

Could you please explain basicly how to apply our model with 26 class or another

>> SEE WHAT'S INSIDE / to this field and I didn't understand anything from the link. I do not

KNOW WHAT TO DO...

I implement your code for my dataset but d1\_loss and d2\_loss converges to 0 and also gan\_loss goes up to 10.



**Adrian Tam** September 11, 2021 at 5:53 am #

With 26 classes, did you tried a 26-bit one-hot encoding?



**CJAY** March 4, 2020 at 6:46 pm #

REPLY ↗

Hey thank you Jason. Very impreesive article.

I'm wondering if CGAN can be used in a regression problem. Since among many GAN, only CGAN have the y label. But I'm not sure how to apply it to non-image problem. For example, I want to generate some synthetic data. I have some real data points of y,x.  $y = f(x_1, x_2, x_3, x_4)$ . y is a non-linear function of  $x_1 \sim x_4$ . I have few hundred  $[x_1, x_2, x_3, x_4, y]$  data. However, I want to have more data since the real data is hard to obtain. So basically I want to generate  $x_1\_fake, x_2\_fake, x_3\_fake, x_4\_fake$ , and  $y\_fake$  where  $y\_fake$  is still the same non-linear function of  $x_1\_fake \sim x_4\_fake$ , i.e.,  $y\_fake = f(x_1\_fake, x_2\_fake, x_3\_fake, x_4\_fake)$ . Is it possible to generate such synthetic dataset using CGAN?

 How to Develop a Pix2Pix GAN for

 Image-to-Image Translation

---

 **Jason Brownlee** March 5, 2020 at 6:31 am #

REPLY ↗

How to Develop a 1D Generative  
Adversarial Network From Scratch in

Yes, you could condition on a numerical variable. A different loss function and activation function would be needed. I recommend experimenting.

 How to Develop a CycleGAN for Image-  
to-Image Translation with Keras

---

 **marpuri ganesh** March 7, 2020 at 2:21 pm #

REPLY ↗

How to Develop a Conditional GAN  
(cGAN) From Scratch

I turned the learning rate, batch size, epochs of the model but no use

 **Jason Brownlee** March 8, 2020 at 6:04 am #

REPLY ↗

from Scratch

Perhaps try some of the suggestions here:

<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

## Loving the Tutorials?

The GANs with Python EBook is

where ~~myplumefo~~ **Really Good stuff** am #

REPLY ↗

>> SEE WHAT'S INSIDE this beginner question.

I'm having trouble understanding some of the syntax when you implement your cGANs

In both define\_discriminator() and define\_generator(), you have paired parenthesis:

examples:

define\_discriminator()

line 6: li = Embedding(n\_classes, 50)(in\_label)

line 9: li = Dense(n\_nodes)(li)

define\_generator()

line 16: gen = Dense(n\_nodes)(in\_lat)

line 17: gen = LeakyReLU(alpha=0.2)(gen)

What is the meaning of the extra (in\_label), (li), (in\_lat), (gen) on the end of each of these lines?

You did not need this in your GANs code.



**Jason Brownlee** March 11, 2020 at 5:29 am #

REPLY ↗

This is the functional API, perhaps start here:

**Never miss a tutorial:**

<https://machinelearningmastery.com/keras-functional-api-deep-learning/>

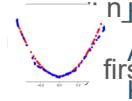


 **Tobias** April 4, 2020 at 2:45 am #

REPLY ↗

Hi Jason!  
[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)  
 visual, a great explanation!

How would you modify the GAN to create n discrete values, i.e. categories, with different classes. Let's say n\_class1=10, n\_class2=15, n\_class=18?

  
[How to Develop a GAN Generative Adversarial Network From Scratch in first thoughts or examples?](#)  
 Keras

Thanks in advance,

 **Tobias** How to Develop a CycleGAN for Image-to-Image Translation with Keras

 **Jason Brownlee** April 4, 2020 at 6:26 am #

REPLY ↗

It would be a much better idea to use a bayesian model instead. This is just a demonstration for how GANs work and a bad example of a generative model for tabular data.  
[Generative Adversarial Network \(WGAN\) From Scratch](#)

 **salman razzaq** April 17, 2020 at 1:10 pm #  
**Loving the Tutorials?**

REPLY ↗

Please guide me how can i modify this code to use it for celebA data set. can i implement the same as that is RGB and there are various labels for a single picture.  
 The [GANs with Python EBook](#) is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE

 **Jason Brownlee** April 17, 2020 at 1:32 pm #

REPLY ↗

This is a common question that I answer here:

[https://machinelearningmastery.com/faq/single-faq/can-you-change-the-code-in-the-tutorial-to-\\_\\_/](https://machinelearningmastery.com/faq/single-faq/can-you-change-the-code-in-the-tutorial-to-__/)

 **Joel** April 19, 2020 at 10:36 am #

REPLY ↗

Hi Jason!

Very nice explanation! Thank you!

I just wonder if there is any pre-trained CGAN or GAN model out there so we can directly use as transfer learning? Specifically, I am interested in Celebra face data.

Thanks again,

 **Never miss a tutorial:**



**Jason Brownlee**

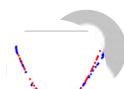
April 19, 2020 at 1:17 pm #

REPLY ↗

**Picked for you** Thanks!

Pre-trained, perhaps. I don't have any sorry.

 How to Develop a Pix2Pix GAN for Image-to-Image Translation



**Nobita Kun** April 20, 2020 at 8:29 pm #

REPLY ↗

Adversarial Network From Scratch in

Thank you Jason, very clear and really useful for a beginner like me.

I have a question about the implementation of the training part.

why should you use half batch for generating real and fake samples but use full\_batch (n\_batch) in

 How to Develop a CycleGAN for Image-to-Image Translation with Keras

`x_real, y_real = generate_real_samples(dataset, half_batch)`

`X_fake, y_fake = generate_fake_samples(g_model, latent_dim, half_batch)`



How to Develop a Conditional GAN (cGAN) From Scratch



**Jason Brownlee** April 21, 2020 at 5:53 am #

REPLY ↗

From Scratch

Thanks.

This is a common heuristic when training GANs, you can learn more here:

<https://machinelearningmastery.com/how-to-train-stable-generative-adversarial-networks/>

The GANs with Python EBook is  
where you'll find the **Really Good** stuff.



21, 2020 at 10:30 am #

REPLY ↗

Thank you for your quick response. I checked the link you gave me, I didn't find information about using half-batch and n-batch. Would you please explain a bit here.



**Jason Brownlee** April 21, 2020 at 11:46 am #

REPLY ↗

From that post:

*Use mini batches of all real or all fake*



**Nobita Kun** April 22, 2020 at 12:49 am #

Thank you very much for your help, Jason.

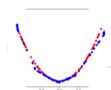
**Never miss a tutorial:****Jason Brownlee** April 22, 2020 at 6:00 am #

You're welcome.

**Picked for you:****pratik korat** July 20, 2020 at 12:49 pm #

REPLY ↗

Image-to-Image Translation

it is functional api that can be used to create branches in your model  
in sequential model you can't do thatHow to Develop a 1D Generative  
Adversarial Network From Scratch in  
Keras**Jason Brownlee** July 20, 2020 at 1:53 pm #

REPLY ↗

**How to Develop a CycleGAN for Image-**  
to-Image Translation with Keras

Agreed.

REPLY ↗

**Akshay** June 21, 2020 at 6:36 pm #  
(cGAN) From Scratch

REPLY ↗

Hi Jason!

Very Nice explanation . Helped me a lot in clarifying my doubts.

I have a request—Can I make same kind Of Explanation for “Context Encoder : Feature learning by  
Generative Adversarial Network (WGAN)  
Inting”.  
From Scratch

Thanks!!

**Loving the Tutorials?**

The GANs with Python EBook is

where you'll find the **Really Good** stuff.

Thanks!

>> SEE WHAT'S INSIDE  
Great suggestion...

REPLY ↗

**Cornelia** July 1, 2020 at 7:49 am #

REPLY ↗

Thanks for the Tutorial, I've been working my way through all the GAN tutorials you provided, it has been super helpful!

I tried training this conditional GAN on different data sets and it worked well. Now I'm trying to train it on a data set with a different number of classes (3 and 5). I changed the n\_classes in every method as well as the label generation after the training and loading of the network of course. However, I get an IndexError and have been unable to solve it. Could you quickly suggest which changes need to be made to change the number of classes in the data set?

Cheers!

**Never miss a tutorial:**

Jason Brownlee July 1, 2020 at 11:18 am #

REPLY ↗

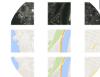
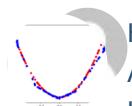


Thanks, happy to hear that.

Sorry to hear that you're having trouble adapting the example.

**Picked for you:**

Perhaps confirm that the number of nodes in the output layer matches the number of classes in the target variable and that the target variable was appropriately one hot encoded.

[How to Develop a Pix2Pix GAN for](#)[Image-to-Image Translation](#)[How to Develop a 3D Generative](#)[Adversarial Network From Scratch in](#)[Keras](#)

REPLY ↗

[How to Develop a Conditional GAN](#)[\(cGAN\) From Scratch](#)

Jason Brownlee July 6, 2020 at 2:09 pm #

REPLY ↗



Interesting idea, you might want to explore using an infogan and have a parameter for each element to Develop a Wasserstein

[Generative Adversarial Network \(WGAN\)](#)[From Scratch](#)

Alex July 6, 2020 at 5:16 pm #

REPLY ↗

**Loving the Tutorials?**

Thank you. I'll read your infogan article.

The [GANs with Python EBook](#) is [\(https://machinelearningmastery.com/how-to-develop-an-information-maximizing-generative-adversarial-network-infogan-in-keras/\)](https://machinelearningmastery.com/how-to-develop-an-information-maximizing-generative-adversarial-network-infogan-in-keras/) where you'll find the [Really Good](#) stuff.

&gt;&gt; SEE WHAT'S INSIDE



Alex July 16, 2020 at 5:28 pm #

REPLY ↗

At present, the model I want to generate is supervised learning, but using InfoGAN should be unsupervised model. Maybe I still use CGAN to generate and control the results?



Jason Brownlee July 17, 2020 at 6:03 am #

REPLY ↗

I recommend using the techniques you think are the most appropriate to your project.



Alex July 16, 2020 at 5:33 pm #

REPLY ↗

**Never miss a tutorial:** For example, I input geometric features x, y, z of various products, then output the process parameters of the product. ex: INPUT (x,y,z) and OUTPUT (temp,pressure,speed),

 Is it possible to predict the process parameters of the product when the geometric characteristics x, y, z of the new product are input?

## Picked for you:

---

 How to Develop a Pix2Pix GAN for Image-to-Image Translation **Jason Brownlee** July 17, 2020 at 6:04 am #

REPLY ↗

It depends on the data, but yes, this is what supervised learning is designed to achieve.

[How to Develop a 1D Generative](#)

[Adversarial Network Frameworks](#) will help:

[Keras](https://machinelearningmastery.com/how-to-define-your-machine-learning-problem/) <https://machinelearningmastery.com/how-to-define-your-machine-learning-problem/>

 How to Develop a CycleGAN for Image-to-Image Translation with Keras **Shaoxuan** July 13, 2020 at 11:30 pm #

REPLY ↗

Hi, Jason,

[How to Develop a Conditional GAN](#)

[\(cGAN\) From Scratch](#) question about the labels in conditional GAN. Instead of several categories, like integers from 0 to can the labels be generated by continuous Uniform distribution(0,1)? So there will be hundreds of labels inputted to the generator or discriminator.

Do you think it is reasonable or doable? Thank you very much!

[How to Develop a Wasserstein](#)

[Generative Adversarial Network \(WGAN\)](#)

[From Scratch](#)



**Jason Brownlee** July 14, 2020 at 6:24 am #

REPLY ↗

Loving the Tutorials?

I don't see why not.

The [GANs with Python](#) EBook is where you'll find the **Really Good** stuff.



>> SEE WHAT'S INSIDE #

REPLY ↗

What changes do I have to make to be able to train with 3 channel images? I changed the input\_shape to (dim, dim, 3) but I still get the error: ValueError: Input 0 of layer conv2d is incompatible with the layer: expected axis -1 of input shape to have value 4 but received input with shape [None, dim, dim, 2]



**Jason Brownlee** July 26, 2020 at 6:13 am #

REPLY ↗

Perhaps start with this model and adapt it to be conditional:

<https://machinelearningmastery.com/how-to-develop-a-generative-adversarial-network-for-a-cifar-10-small-object-photographs-from-scratch/>

 **Richard Macwan** August 27, 2020 at 4:36 am #

REPLY ↗

**Never miss a tutorial!**

I had to change the Embedding layer's dimensions for the conditional Gan otherwise tf complained that it could not reshape (32,1,50) to (32,7,7).

 I changed the 50 to 49 as li = Embedding(n\_classes,49)(in\_label). Am I missing something or it is a typo?

**Picked for you:**

 How to [Develop a CycleGAN for Image-to-Image Translation](#) Jason Brownlee August 27, 2020 at 6:25 am #

REPLY ↗

That is odd, sorry to hear that.

Did you copy all other code examples as-is without modification?

[How to Develop a 1D Generative](#)

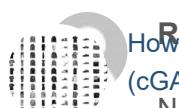
Are your libraries up to date?

[Adversarial Network From Scratch in Keras](#)

Do these tips help:

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me/>

 [How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

 How to [Develop a Conditional GAN \(cGAN\) From Scratch](#) Richard Macwan August 27, 2020 at 8:56 pm #

REPLY ↗

No since I haven't purchased the book, I don't have the code. I did find the problem though, and

I was silly to ask the question! It was obvious that I had missed adding a Dense layer before Reshaping

the generator!

[How to Develop a Wasserstein](#)

[Generative Adversarial Network \(WGAN\)](#)

Thanks for your prompt reply.

From Scratch

 Loving the Tutorials? Jason Brownlee August 28, 2020 at 6:42 am #

REPLY ↗

The GANs with Python EBook is  
I'm happy to hear that you solved your problem!  
where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE



Nir Regev September 2, 2020 at 12:59 am #

REPLY ↗

Hi Jason,  
great post (again).

How would the conditional DCGAN would change if instead of label (condition) input, I have a facial landmark image of some dimension (e.g. a grayscale of 28,28) ? The generator, in this case, should generate an image that corresponds to the landmark image, and so is the discriminator should "judge" the image according to the landmark condition.

Specifically I'm struggling to understand what should be in the in\_label and li in the define\_discriminator method. thanks

 Jason Brownlee September 2, 2020 at 6:31 am #

REPLY ↗

Not dramatically, perhaps just adjustments to the model for the new input shape.

**Never miss a tutorial:**



Omar September 22, 2020 at 1:15 am #

REPLY ↗

**Picked for you:**

Hello,

Thanks a lot for this tutorial! I really needed this.

[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)

one question, What are the changes needed in the generator and the discriminator if I am using a custom dataset, with dimensions (64,64,3) not (28,28,1)?

This is really bugging me, I know it's simple but I think I might be overseeing something.

[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

I also found your help very useful!

[How to Develop a CycleGAN for Image-to-Image Translation With Keras](#)

September 22, 2020 at 6:51 am #

REPLY ↗

Change the expected input shape for the discriminator and perhaps add more layers to support the larger images.

[How to Develop a Conditional GAN](#)

Add more layers to the generator to ensure the output shape matches the expected size.

How many layers and what types – you will have to discover the answer via trial and error.

[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)

Omar September 22, 2020 at 6:43 pm #

REPLY ↗

Thanks, will have a look into this.

**Loving the Tutorials?**

One more thing, my dataset is loaded as a tf batch dataset (using keras image dataset from directory). How can adjust the code to train this instead of the MNIST fashion?

The [GANs with Python](#) eBook is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE



Jason Brownlee September 23, 2020 at 6:35 am #

REPLY ↗

Sorry, I don't know about "tf batch dataset".



Kim Quinto October 2, 2020 at 9:55 pm #

REPLY ↗

Hello,

Thanks for the detailed blogs! I am trying to improve a classifier that was trained on a small imbalance dataset, and I am considering using CGAN to extend my dataset and making it a balanced one, but here comes the question: Can I use the whole dataset (train, validation and test) to train my CGAN and then use the generated images to extend and balance my classifier? or should I use the training set only? I am a little bit confused if using the whole dataset is completely fine since the generated images will be from a different distribution. I couldn't find any answer for my confusion, so what do you think?

**Never miss a tutorial:**

Kim Quinto October 2, 2020 at 9:58 pm #

REPLY ↗



I could reuse in that case, my classifier will be validated and tested on generated images that were trained on these datasets to be generated.

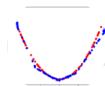
**Picked for you:**

How to Develop a Pix2Pix GAN for Image-to-Image Translation

Tiwahade Usman December 25, 2020 at 1:57 am #

REPLY ↗

You apply it only to your training data



How to Develop a 1D Generative Adversarial Network From Scratch in Keras

Jason Brownlee October 3, 2020 at 6:07 am #

REPLY ↗



How to Develop a CycleGAN for Image-to-Image Translation with Keras

No, I think it would be valid to only the training dataset to train the GAN as part of your experiment.



How to Develop a Conditional GAN (cGAN) From Scratch

Sang Young Lim October 6, 2020 at 1:19 am #

REPLY ↗

In regard of your nice post below,  
[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\)](https://machinelearningmastery.com/generative-adversarial-network-loss-functions/)

...  
**Loving the Tutorials?**

Because I think your code and explanations imply following statement  
 "In practice, this is also implemented as a binary classification problem, like the discriminator. Instead of maximizing the loss, we fit the labels for real and fake images and minimize the cross-entropy."

I am trying to use Bayesian Optimization method on this cDCGAN above, and I got lost deciding to define the evaluation function to look for bigger g\_loss or smaller g\_loss (on average).

I have done 500 epoch on this code above, but could not find g\_loss is going down or up.  
 Thank >> SEE WHAT'S INSIDE Great work!



Jason Brownlee October 6, 2020 at 6:57 am #

REPLY ↗

In general, for GANs, no:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-gan-not-converging>



Eoin Kenny October 7, 2020 at 8:39 pm #

REPLY ↗

Hi Jason thanks for this.

Quick question, does the embedding work with floating point numbers? I don't want to only have ints for input here, I want a float too, is that possible?

Thanks you.  
**Never miss a tutorial:**



Jason Brownlee

October 8, 2020 at 8:30 am #

REPLY ↗

Picked for you! Embedding layers in general? Yes.

How to Develop a Pix2Pix GAN for Image-to-Image Translation

REPLY ↗

Harshi Rao November 22, 2020 at 1:07 am #

Hi Jason,  
 I have been following your blogs, posts and newsletters for the past few years!  
 Do you have any advice on how to apply GANs for document generation?

Thanks in advance.

How to Develop a CycleGAN for Image-to-Image Translation with Keras

REPLY ↗

Jason Brownlee November 22, 2020 at 6:56 am #

How to Develop a Conditional GAN (cGAN)  
 Thanks Scratch

I would recommend "language models" for document generation, not GANs:

[https://machinelearningmastery.com/?s=language+models&post\\_type=post&submit=Search](https://machinelearningmastery.com/?s=language+models&post_type=post&submit=Search)  
 How to Develop a Wasserstein

Generative Adversarial Network (WGAN)  
 From Scratch

REPLY ↗



Ali November 23, 2020 at 8:45 pm #

Loving the Tutorials?

Hi Jason,

My question is naive, but would appreciate if you answer it.

The [GANs with Python EBook](#) is

Assume that I have 1D data and want to have only dense layers in both models. Here is the code for my disc model definition:

```
def def    >> SEE WHAT'S INSIDE  =(10,1), n_classes=8):
# label input
in_label = Input(shape=(1,))

li = Embedding(n_classes, 50)(in_label)

n_nodes = in_shape[0]
li = Dense(n_nodes)(li)
# reshape to additional channel
li = Reshape((n_nodes, 1))(li)

in_data = Input(shape=in_shape)
# concat label as a channel
merge = Concatenate()([in_data, li])
hidden1 = Dense(64, activation='relu')(merge)
hidden2 = Dense(64, activation='relu')(hidden1)
# output
out_layer = Dense(1, activation='sigmoid')(hidden2)
# define model
```

```
model = Model([in_data, in_label], out_layer)
```

**Never miss a tutorial:**

# compile model

```
opt = Adam(lr=0.0002, beta_1=0.5)
```

```
model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
```

return model

**Picked for you:**

The problem is that when I use d\_model.predict(), the output has 3 dimensions instead of 2. In fact, the

[we should Develop a Pix2Pix GAN for 10, 1](#) where 10 is the input dimension. Please let me know what I

[missing here](#)

[Image Translation](#)

---

 How to Develop a 1D Generative

**Jason Brownlee** November 24, 2020 at 6:19 am #

REPLY ↗

Keras

Perhaps review a plot or summary of the model to confirm it was constructed the way you intended.

 [How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

 **Tiwalade Usman** December 25, 2020 at 2:00 am #

[How to Develop a Conditional GAN](#)

[\(cGAN\) From Scratch](#)

[How do I apply cGAN to augment images with multi-labels?](#)

REPLY ↗

---

 [How to Develop a Wasserstein](#)

[Generative Adversarial Network \(WGAN\)](#)

[From Scratch](#)

REPLY ↗



Good question. I have not explored this, perhaps use a little trial and error to see what works well.

**Loving the Tutorials?**

Let me know how you go.

The [GANs with Python](#) EBook is  
where you'll find the **Really Good** stuff.



>> SEE WHAT'S INSIDE December 26, 2020 at 8:44 pm #

REPLY ↗

Hi Jason,

I am interested to know if it is possible to checkpoint your gan model. Since typical metrics e.g. loss and accuracy do not work for GANs, we have to define custom metrics. Using keras, we can easily do this for a classifier like the disc model, but I don't know how to do this with the gen model. Please assume that we have a metric to analyze the quality of the generated images.

---



**Jason Brownlee** December 27, 2020 at 5:00 am #

REPLY ↗

Not really as loss does not relate to image quality.

You can save after each manually executed epoch if you like.

**Never miss a tutorial:**

Nadjib December 28, 2020 at 12:06 pm #

REPLY ↗



Hi, I like your the great tutorial. Is there a way of using more than one condition vector ? could we use multiple condition vectors then concatenate them ?

**Picked for you:**

How to Develop a Pix2Pix GAN for  
Jason Brownlee December 28, 2020 at 1:16 pm #

REPLY ↗

You're welcome.



How to Develop a 1D Generative Adversarial Network From Scratch in  
Keras



How to Develop a CycleGAN for Image-to-Image Translation with Keras

REPLY ↗

Thanks for teaching me ML. 2 Questions:

I notice train\_on\_batch is passed a half\_batch real & half\_batch fake data. Are weights somehow only updated on loading a full batch?

[How to Develop a Conditional GAN with GANs From Scratch](#)

Would it be such a high 50d Embedding for mapping only 10 classes?

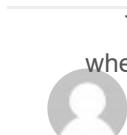


How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch

REPLY ↗

Weights are updated after each batch update.

50d embedding is common, you can try alternatives.

**Loving the Tutorials?**

The GANs with Python EBook is  
where you'll find the **Really Good** stuff.

Gili January 5, 2021 at 8:45 pm #

REPLY ↗

>> SEE WHAT'S INSIDE great tutorial!

I'm trying to create a conditional gan for time series data so my model is using LSTMs instead of CNNs. I'm having trouble understanding how to reshape my input and the labels embeddings. In my case, instead of a 28x28 image, I have a time-series sample of shape: time\_steps, n\_features As you know, an LSTM needs the input to be [n\_samples, time\_steps, n\_features] but now I also need to add the labels and I will get 4 dimensions instead of 3. do you have any suggestions on how to do this right? thanks a lot!



Jason Brownlee January 6, 2021 at 6:28 am #

REPLY ↗

Perhaps this will help:

<https://machinelearningmastery.com/faq/single-faq/what-is-the-difference-between-samples-timesteps-and-features-for-lstm-input>

**Never miss a tutorial:**

Dennis February 16, 2021 at 3:51 pm #

REPLY ↗



Thanks for your great post.

In you code when generate fake image;

**Picked for you:**

```

1 def generate_fake_samples(generator, latent_dim, n_samples):
2     # generate points in latent space
3     z_input, labels_input = generate_latent_points(latent_dim, n_samples)
4     # predict outputs
5     images = generator.predict([z_input, labels_input])
6     # create class labels
7     y = zeros((n_samples, 1))
8     return [images, labels_input], y

```

Adversarial Network From Scratch in Keras

But based on the paper, I see other people use real label inputs(same as from real image sample) and `z_input` to generator fake image, then test d model weights based on this fake image. But it seems this won't influence the model. I want to ask which one is correct? Is there any difference?

[How to Develop a CycleGAN for image-to-Image Translation with Keras](#)

Thank you!

**How to Develop a Conditional GAN**

(cGAN) From Scratch

Jason Brownlee February 17, 2021 at 5:25 am #

REPLY ↗

There are many different types of models.

[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\)](#)

From Scratch

**Sree Niranjan** March 4, 2021 at 5:24 pm #

REPLY ↗

Hi Jason, The GANs with Python EBook is

I am using able to do real GAN for my time-series data generation. Based on my application, the generator loss converges at certain point and then it increases. How can I include the early stopping in my model? >> SEE WHAT'S INSIDE > below certain threshold, the training should get terminated and the model be saved.

Jason Brownlee March 25, 2021 at 4:41 am #

REPLY ↗

Generally GANs do not converge:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-gan-not-converging>

David Yen March 27, 2021 at 2:23 pm #

REPLY ↗

Hi Jason, I am trying to apply your conditional GAN code to a CT scan dataset with 256×256 input images. I added a few more layers in discriminator and generator models. A sample of generator code changes I made is shown below. The training from this modified code takes hours without showing any errors or results. Any idea what went wrong? Thanks.

# linear multiplication  
**Never miss a tutorial:**

n\_nodes = 4 \* 4



li = Dense(n\_nodes)(li)

# image generator input

in\_lat = Input(shape=(latent\_dim,))

How to Develop a Pix2Pix GAN for Image Translation

How to Develop a 1D Generative Adversarial Network From Scratch in Keras

gen = Dense(n\_nodes)(in\_lat)  
 gen = LeakyReLU(alpha=0.2)(gen)  
 = Reshape((4, 4, 1024))(gen)  
 Adversarial Network From Scratch in Keras

merge = Concatenate()([gen, li])

# upsample to 8x8

How to Develop a cGAN for Image Translation with Keras

# upsample to 16x16

gen = Conv2DTranspose(256, (5,5), strides=(2,2), padding='same')(gen)

How to Develop a Conditional GAN (cGAN) From Scratch

gen = Conv2DTranspose(128, (5,5), strides=(2,2), padding='same')(gen)

aen = LeakyReLU(alpha=0.2)(gen)

= How to Develop a Wasserstein

Generative Adversarial Network (WGAN) From Scratch

gen = LeakyReLU(alpha=0.2)(gen)

# upsample to 128x128

gen = Conv2DTranspose(32, (5,5), strides=(2,2), padding='same')(gen)

Loving the Tutorials?

gen = LeakyReLU(alpha=0.2)(gen)

# upsample to 256x256

The GANS with Python EBook is

gen = Conv2DTranspose(16, (5,5), strides=(2,2), padding='same')(gen)

gen = LeakyReLU(alpha=0.2)(gen)

# output >> SEE WHAT'S INSIDE

out\_layer = Conv2D(1, (256,256), activation='tanh', padding='same')(gen)



**Jason Brownlee** March 29, 2021 at 6:01 am #

REPLY ↗

Perhaps some of these suggestions will help:

<https://machinelearningmastery.com/faq/single-faq/how-do-i-speed-up-the-training-of-my-model>



**Jessie** April 19, 2021 at 10:24 pm #

REPLY ↗

Hi Jason, I have question about load model and continue training. When I end the set epoch, I save the g\_model to h5. If I want to use the model for continue training, which model I need. Need to save the d\_model and gan model? Thanks.

**Never miss a tutorial:**

Jason Brownlee April 20, 2021 at 5:58 am #

REPLY ↗



spe You will need to save/load the G and D and composite models.

**Picked for you:**

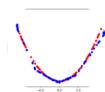
How to Develop a Pix2Pix GAN for

REPLY ↗



Image-to-Image Translation

It help me to continue training. Thanks you for reply.



How to Develop a 1D Generative

Adversarial Network From Scratch in

Keras Jason Brownlee April 21, 2021 at 5:51 am #

REPLY ↗



You're welcome.

How to Develop a CycleGAN for Image-  
to-Image Translation with Keras

How to Develop a Conditional GAN

REPLY ↗

(cGAN) From Scratch Thanks for your reply. I success to do continue training.

How to Develop a Wasserstein  
Generative Adversarial Network (WGAN)  
From Scratch Omnia April 28, 2021 at 1:19 pm #

REPLY ↗

Thanks a lot for all of your works, Jason. Great.

Why didn't save gan\_model and load that one!?

I try to do it myself, so in mine, but couldn't!

May you load gan\_model and test?

The [GANs with Python](#) EBook iswhere you'll find the **Really Good** stuff.

&gt;&gt; SEE WHAT'S INSIDE

April 29, 2021 at 6:22 am #

REPLY ↗

You're welcome.

The composite model is only used during training, then discarded.



xin May 6, 2021 at 2:07 am #

REPLY ↗

I used a dataset from (<http://vision.ucsd.edu/~leekc/ExtYaleDatabase/ExtYaleB.html>) and wanted to generate face images with Conditional gan, I referenced the network architecture from ([https://github.com/GANs-in-Action/gans-in-action/blob/master/chapter-8/Chapter\\_8\\_CGAN.ipynb](https://github.com/GANs-in-Action/gans-in-action/blob/master/chapter-8/Chapter_8_CGAN.ipynb)) and (<https://github.com/mvedang/Conditional-GAN-CIFAR-10>) network architecture, but no matter how many times I train it, the resulting image is still noise.

Even if I adjust the parameters of the Discriminator and Generator such as layers, neurons, etc., the generated results are still noise.

Can you help me to correct the Generator and Discriminator models I have constructed? code and result  
**Never miss a tutorial:**  
[link:https://colab.research.google.com/drive/14ul6BeXpVnIO3TVmfvSiR0kkYUuWfvn-?usp=sharing](https://colab.research.google.com/drive/14ul6BeXpVnIO3TVmfvSiR0kkYUuWfvn-?usp=sharing)



```
def build_cgan_generator(z_dim):
```

```
    z_input = Input(shape=(z_dim,))
```

```
    label_input = Input(shape=(1,), dtype='int32')
```

```
    label_embedding = Embedding(classes, output_dim=z_dim, input_length=1)(label_input)
```

```
    label_embedding = Flatten()(label_embedding)
```

```
    join_represent = Multiply()([z_input, label_embedding])
```

```
    x = Dense(4*4*256)(join_represent)
```

```
    x = Reshape((4, 4, 256))(x)
```

```
    x = Conv2DTranspose(64, kernel_size=3, padding='same', strides=2)(x)
```

```
#x = BatchNormalization(momentum=0.8)(x)
```

```
x = LeakyReLU(0.01)(x)
```

```
    x = Conv2DTranspose(128, kernel_size=3, padding='same', strides=2)(x)
```

```
#x = BatchNormalization(momentum=0.8)(x)
```

```
x = LeakyReLU(0.01)(x)
```

```
    x = Conv2DTranspose(128, kernel_size=3, padding='same', strides=2)(x)
```

```
#x = BatchNormalization(momentum=0.8)(x)
```

```
x = LeakyReLU(0.01)(x)
```

```
    x = Conv2DTranspose(64, kernel_size=3, padding='same', strides=2)(x)
```

```
#x = BatchNormalization(momentum=0.8)(x)
```

```
x = LeakyReLU(0.01)(x)
```

```
    x = Conv2DTranspose(32, kernel_size=3, padding='same', strides=2)(x)
```

```
#x = BatchNormalization(momentum=0.8)(x)
```

```
x = LeakyReLU(0.01)(x)
```

```
x = Conv2DTranspose(3, kernel_size=4, padding='same', strides=2)(x)
```

## Loving the Tutorials?

```
output = keras.layers.Activation('tanh')(x)
```

The [GANs with Python](#) EBook is

model = keras.Model([z\_input, label\_input], output)

where you'll find the [Really Good](#) stuff.

```
tf.keras.utils.plot_model(model, to_file='generator.png', show_shapes=True)
```

```
return ! >> SEE WHAT'S INSIDE
```

Discriminator.

```
def build_cgan_discriminator(img_shape):
```

```
    img_input = Input(shape=img_shape)
```

```
    label_input = Input(shape=(1,))
```

```
    label_embedding =
```

```
    Embedding(classes, output_dim=np.prod((img_shape[0], img_shape[1], 1)), input_length=1)(label_input)
```

```
    label_embedding = Flatten()(label_embedding)
```

```
    label_embedding = Reshape((img_shape[0], img_shape[1], 1))(label_embedding)
```

```
    concatenated = Concatenate(axis=-1)([img_input, label_embedding])
```

```
x = layers.Conv2D(64, kernel_size=3, strides=2, padding='same')(concatenated)
```

```
x = LeakyReLU(0.01)(x)
```

```
x = layers.Conv2D(64, kernel_size=3, strides=2, padding='same')(x)
```

```
#x = BatchNormalization(momentum=0.8)(x)
```

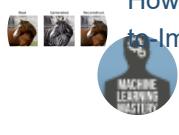
```
x = LeakyReLU(0.01)(x)
```

x=layers.Conv2D(128,kernel\_size=3,strides=2,padding='same')(x)  
**Never miss a tutorial:**  
 #x = BatchNormalization(momentum=0.8)(x)

x=LeakyReLU(0.01)(x)  
 x=layers.Conv2D(64,kernel\_size=3,strides=2,padding='same')(x)  
 #x = BatchNormalization(momentum=0.8)(x)  
**Picked for you:**  
 X=LeakyReLU(0.01)(x)

  
**National Day** How to Develop a Pix2Pix GAN for Image Translation (age) Transformation  
 outputs = layers.Dense(1,activation='sigmoid',name='Output')(x)

model = keras.Model([img\_input,label\_input],outputs)  
**How to Develop a 1D Generative Adversarial Network From Scratch in Keras**  
 as.utils.plot\_model(model,to\_file='discriminator.png',show\_shapes=True)  
 return model

  
**How to Develop a CycleGAN for Image-to-Image Translation with Keras**  
**Jason Brownlee** May 6, 2021 at 5:48 am #

REPLY ↗

Sorry, I don't have the capacity to review/debug your code. Perhaps these tips will help:

**How to Develop a Conditional GAN (cGAN) From Scratch**  
<https://machinelearningmastery.com/faq/single-faq/can-you-read-review-or-debug-my-code>

  
**How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch**  
**Omid** May 14, 2021 at 7:22 pm #

REPLY ↗

Thanks a lot for all.

I implemented cGAN on a different data. But the loss of generator is increasing and accuracy for both generator and discriminator are 100% in different epochs and batch size.

**Loving the Tutorials?**

The **GANs with Python** EBook is where you'll find the **Really Good** stuff.

  
 >> SEE WHAT'S INSIDE y 15, 2021 at 6:30 am #

REPLY ↗

Generally, loss is not a good indicator for GAN performance:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-gan-not-converging>

Perhaps try running the example a few times?

Perhaps try adapting the model configuration?

  
**nadjoua** June 22, 2021 at 4:14 am #

REPLY ↗

hi sear,  
 how can we use cgan for dataset celeba?

  
**Jason Brownlee** June 22, 2021 at 6:33 am #

REPLY ↗

Not sure it is appropriate.  
**Never miss a tutorial:**



**nadjoua** July 31, 2021 at 9:22 pm #

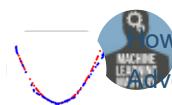
REPLY ↗

### Picked for you:

Please sear, can you help me; how can i modify this code to use it for celebA data set. can i implement the same as that is RGB and there are various labels for a single picture.

[How to Develop a Pix2Pix GAN for](#)

[Image-to-Image Translation](#)



**Jason Brownlee** August 1, 2021 at 4:51 am #

REPLY ↗

Adversarial Network From Scratch in Keras Perhaps this tutorial will help:

<https://machinelearningmastery.com/how-to-interpolate-and-perform-vector-arithmetic-with-faces-using-a-generative-adversarial-network/>

[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)



**nadjoua** August 1, 2021 at 10:23 pm #

REPLY ↗

[How to Develop a Conditional GAN \(cGAN\) From Scratch](#)

thank you so much sear



[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\)](#)

**Jason Brownlee** August 2, 2021 at 4:53 am #

REPLY ↗

You're welcome.

### Loving the Tutorials?

The GANs with Python EBook is  
**Tim** August 17, 2021 at 5:57 pm #

REPLY ↗

where you'll find the **Really Good** stuff.

... >> SEE WHAT'S INSIDE ...

e a great help. So I had a silly idea take this example and generate 60000 to train a Fasion Mnist classification network from your other tutorial, and to my astonishment the generated dataset gave much higher accuracy on a validation set than the real training dataset. I didn't change anything in the example but the training set, and the result was over 99% accuracy, compared to around 90% in 10 epochs for the real one. This seems too good to be true considering that the best method from the benchmark gets ~96% (<https://paperswithcode.com/sota/image-classification-on-fashion-mnist>), so I would want to ask if you have any idea why this might have happened?



**Adrian Tam** August 18, 2021 at 3:04 am #

REPLY ↗

That might mean your fake samples are not fake enough. This is a difficult problem to solve. But think in this way, if your fake samples are too simple to identify and use it to train, the machine will not learn anything useful. Just like giving you unchallenging exercises to do, even doing a lot, you learned nothing.

**Never miss a tutorial:**

Tim August 18, 2021 at 5:36 am #

REPLY ↗



Hi Alin, thanks for reply. But if the fake samples weren't good enough to substitute the training set, then I would expect the accuracy on the validation set (a portion dataset that GAN generator has never seen) to be much lower. What I'm wondering is why a simple classification network separate to the GAN trained with only fake 60k samples yielded higher accuracy than when trained with the official 60k samples, I imagined it would be the other way around.

**Picked for you:**

[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)

[How to Develop a 1D Generative](#)[Adversarial Network From Scratch](#)

Keras

Terry Taerum September 18, 2021 at 11:51 am #

REPLY ↗



Without looking at the data you generated, I cannot really tell. But one possibility to explain is this: If the original sample is a bigger problem (e.g. identify a thousand objects) How to Develop a CycleGAN for Image-to-Image Translation with Keras and the generated samples is only a smaller problem (e.g. identify a dog vs a car), then you will likely see that the accuracy of latter is better than former. One way to prove you are actually not doing any better is to use the generated samples to train your network and use the official samples for validation.

[How to Develop a Conditional GAN](#)[\(cGAN\) From Scratch](#)

Hope this helps!

[How to Develop a Wasserstein](#)[Generative Adversarial Network \(WGAN\)](#)

From Scratch

Terry Taerum September 13, 2021 at 11:31 am #

REPLY ↗

I've come back to this piece of brilliant work on your part.

What I don't understand is how the information from the d\_model gets transferred to the g\_model via the gan\_model so quickly. I continue to make modifications and the following is a cryptic piece within the training loop. The GANs with Python Example reused some calculations.

```
where you'll find the Really Good stuff.
[X_real, labels_real], y_real = generate_real_samples(dataset, n_batch)
d_learr >> SEE WHAT'S INSIDE    lr/g_loss, counter)
K.set_value(gan_model.optimizer.learning_rate, d_learning_rate)
d_loss1, _ = d_model.train_on_batch([X_real, labels_real], y_real)
z_input, labels = generate_latent_points(latent_dim, n_batch)
X_fake = g_model.predict([z_input, labels])
y_fake = zeros((n_batch, 1))
d_loss2, _ = d_model.train_on_batch([X_fake, labels], y_fake)
y_gan = ones((n_batch, 1))
g_learning_rate = calculate_learning_rate(lr, counter)
K.set_value(gan_model.optimizer.learning_rate, g_learning_rate)
g_loss = gan_model.train_on_batch([z_input, labels], y_gan)
if (j%10==0):
print('>%d/%d, %d/%d, d1=% .3f, d2=% .3f g=% .3f %
(i+1,n_epochs, j+1, bat_per_epo, d_loss1, d_loss2, g_loss))
```

You can see I don't recalculate the latent\_points within each loop. It makes little difference in execution time but it illustrates I only need to calculate my latent points at the beginning of each loop and re-use them. All this only deepens the mystery for me. As I understand it, the g\_model weights get updated

when they are passed through the gan\_model (which is the only place they could be updated). So much of this appears to happen in the background.

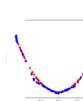
**Never miss a tutorial:**  also detect when learning rate and pull back the rate on d\_model when gan\_model appears to be heading towards a collapse.

**Picked for you:** def calculate\_learning\_rate(lr, counter):

```
lr = lr * 0.998
```

 Learn How to Develop a Pix2Pix GAN for Image-to-Image Translation manipulating the weights.

Hopefully I'm not doing anything too far off the charts.

 [How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

 **Liangliang Xiang** October 6, 2021 at 8:31 pm #

[How to Develop a CycleGAN for Image-](#)

REPLY ↗

 [to-Image Translation with Keras](#)  
Hi Jason, Could this conditional GAN structure be used for the 1D inertial sensor data generation?

 [How to Develop a Conditional GAN \(cGAN\) From Scratch](#)

 **Adrian Tam** October 7, 2021 at 3:00 am #

REPLY ↗

 [How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)

 **Liangliang Xiang** October 16, 2021 at 1:02 pm #  
**Loving the Tutorials?**

REPLY ↗

Because most cases are utilized for images generation rather than 1d sensor data. So, The [GANs with Python EBook](#) is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE



**Adrian Tam** October 20, 2021 at 8:44 am #

REPLY ↗

I believe it should work but I can't see your data. Always the only way to confirm is to experiment!



**Red1** November 18, 2021 at 11:46 pm #

REPLY ↗

Thanks a lot for all.

i would like to implemente this code to generate data (data augmentation) in order to balcance data data and then make classification

the probleme how can i modify this code to use it for credit cardt fraud detection data set to generate fraudulent trabsaction classe 1 in .

**Never miss a tutorial:**

Ayat Firas December 30, 2021 at 5:34 am #

REPLY ↗



Thank you, or to receive information. Is it possible to give me the dataset for this code?

Thank you very much

**Picked for you:**

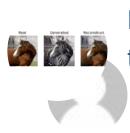
How to Develop a Pix2Pix GAN for Image-to-Image Translation **James Carmichael** December 30, 2021 at 10:04 am #

REPLY ↗

Hello Ayat... You may find the datasets at the following location:



<http://yann.lecun.com/exdb/mnist/>  
Adversarial Network From Scratch in  
Regards,  
Keras



How to Develop a CycleGAN for Image-to-Image Translation with Keras **Harry** January 24, 2022 at 5:59 pm #

REPLY ↗

Hello sir,

[How to Develop a Conditional GAN \(cGAN\) From Scratch](#) blog has helped me a lot in learning the basics of C-GAN. I am applying C-GAN for a physics tel to learn the phase transition in theory.

If I want to generate data (say image pixels in your model) for intermediate labels (not part of training)

then how to do that?



[How to Develop a Wasserstein](#)

[Generative Adversarial Network \(WGAN\)](#)

From Scratch



Loving the Tutorials? **James Carmichael** January 25, 2022 at 10:48 am #

REPLY ↗

The [GANs with Python Book](#) is for the feedback! You may find the following of interest:  
where you'll find the **Really Good** stuff.  
<https://www.toptal.com/machine-learning/generative-adversarial-networks>

>> SEE WHAT'S INSIDE



**Shubh** January 28, 2022 at 9:30 pm #

REPLY ↗

Dear sir,

I have 10 class labels but all of them are in range (0,1) like [0.20169, 0.22169, ....] so what number of n\_classes shall I take in above code? I tried with n\_class=10 but then d\_loss became zero and I didn't get expected output



**James Carmichael** February 18, 2022 at 12:57 pm #

REPLY ↗

Hi Shubh,

You may be working on a regression problem and achieve zero prediction errors.

Alternately, you may be working on a classification problem and achieve 100% accuracy.

This is unusual and there are many possible reasons for this, including:  
**Never miss a tutorial:**

You are evaluating model performance on the training set by accident.

 Your hold out dataset (train or validation) is too small or unrepresentative.

You have introduced a bug into your code and it is doing something different from what you expect.

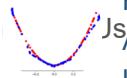
Your prediction problem is easy or trivial and may not require machine learning.

**Picked for you:**  
 The most common reason is that your hold out dataset is too small or not representative of the broader problem.

 [How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)

Using k-fold cross-validation to estimate model performance instead of a train/test split.

Gather more data.

 [How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)

Use a different split of data for train and test, such as 50/50.

Keras

REPLY ↗

 **Shubh** May 28, 2022 at 10:39 am #  
[How to Develop a CycleGAN for Image-to-Image Translation with Keras](#)

Dear sir,

I have modified above code for my use but I am facing a problem i.e, I have 10 labels which are real numbers between 0 and 1 so when I am using above code I get an error because of embedding layer,  
 [How to Develop a Conditional GAN From Scratch](#)

REPLY ↗

 **James Carmichael** January 31, 2022 at 11:04 am #  
[How to Develop a Wasserstein Generative Adversarial Network \(WGAN\) From Scratch](#)

Hi Shubh...What is the exact error you are encountering so that I may better assist you? In general I cannot debug your code, however something may stand out immediately if you can provide the exact error message (pref).

**Loving the Tutorials?**

The [GANs with Python](#) EBook is where you'll find the **Really Good** stuff.

REPLY ↗

 **Adrian** May 22, 2022 at 5:01 pm #  
 >> SEE WHAT'S INSIDE

Hello sir,

Thanks a lot for this blog, it helped me so much. I have a problem when I try to train the gan. I adapted the code to my particular case (images of 6x6), but the output of the train phase is clearly different from yours. Apparently I've only changed the input shape and the output (the upsample and down sample are modified too), but nothing more.

My output is:

```
>1, 1/20, d_loss_real=18834628.000, d_loss_fake=0.696 g_loss=0.690
>1, 2/20, d_loss_real=915644.875, d_loss_fake=0.701 g_loss=0.686
>1, 3/20, d_loss_real=35437.840, d_loss_fake=0.706 g_loss=0.681
>1, 4/20, d_loss_real=0.000, d_loss_fake=0.713 g_loss=0.676
>1, 5/20, d_loss_real=0.000, d_loss_fake=0.719 g_loss=0.669
...
>100, 18/20, d_loss_real=0.000, d_loss_fake=0.023 g_loss=3.808
>100, 19/20, d_loss_real=0.000, d_loss_fake=0.022 g_loss=3.819
>100, 20/20, d_loss_real=0.000, d_loss_fake=0.024 g_loss=3.763
```

**WARNING:tensorflow:** Compiled the loaded model, but the compiled metrics have yet to be built.  
**Never miss a tutorial:** `model.compile_metrics` will be empty until you train or evaluate the model.



## Picked for you:



**James Carmichael** May 24, 2022 at 10:10 am #

REPLY ↗



How to Develop a Pix2Pix GAN for

Hi Adrian... The following discussion may be of interest to you:



<https://stackoverflow.com/questions/67970389/warningtensorflowcompiled-the-loaded-model-but-the-compiled-metrics-have-yet>

How to Develop a 1D Generative



<https://github.com/theAIbytesCode/yolov4-deepsort/issues/79>

Keras



How to Develop a CycleGAN for Image-to-Image Translation with Keras

REPLY ↗



Hi please tell me How to save images after Images Generation through GAN, Wait for your



How to Develop a Conditional GAN (cGAN) From Scratch



**James Carmichael** June 18, 2022 at 10:44 am #

REPLY ↗

Generative Adversarial Network (WGAN)

Hi Ahtisham... The following may be of interest to you:

From Scratch

<https://stackoverflow.com/questions/71452209/save-gan-generated-images-one-by-one>

## Loving the Tutorials?



The [GANs with Python EBook](#) is where you'll find the **Really Good** stuff.

REPLY ↗

Hello Sir

>> SEE WHAT'S INSIDE

How can I use images that have a line of text with a different shape in width and height in my dataset such as h=200, w= 2048 and not n×n in as cGan input, I tried using resize, but it got distorted.



**James Carmichael** August 13, 2022 at 5:59 am #

REPLY ↗

Hi M.M... You may find the following resource of interest:

<https://machinelearningmastery.com/how-to-load-convert-and-save-images-with-the-keras-api/>



**M.M** August 15, 2022 at 8:49 pm #

REPLY ↗

Hi, Thank you for replying.

Q/ I am working on a topic (Data Augmentation by using cGAN), for Arabic text handwriting images. The dataset contains images of large sizes and different dimensions (M\*N).

I tried to use the code sent by you in order to convert the images to square sizes (N\*N), but the result, the images are unclear for reading, can we use this unclear image in cGAN and give a result? To use them in next step in recognition of text ?



– Is there a better way to prevent image distortion when converted it to (N\*N),

or

**Picked for you:**

– can i use an image of different sizes in cGAN as input instead of resizing it?

How to Develop a Pix2Pix GAN for Image-to-Image Translation

**Dee** September 14, 2022 at 12:47 am #

REPLY ↗

How to Develop a 1D Generative Adversarial Network From Scratch in Keras  
My thanks for this cGAN's tutorial!

For cGAN, it can accept label (via one-hot encoding) to generate targeted images with the corresponding object. Could you please let me know if it is possible to train a cGAN model to generate images with multiple (e.g., two or three) targeted objects? For example, if we want to generate images containing 'Dress', 'Shirt' and 'Bag', so the input label vector could be [0, 0, 0, 1, 0, 0, 1, 0, 1, 0] (use '1' to activate the object that we want). I am just wondering if this is doable for the cGAN model?

How to Develop a Conditional GAN (cGAN) From Scratch

How to Develop a Wasserstein Generative Adversarial Network (WGAN) **James Carmichael** September 14, 2022 at 6:01 am #

REPLY ↗

From Scratch

Hi Dee... You are very welcome! I would highly recommend the following resources to support your understanding of GANs:

<https://machinelearningmastery.com/tour-of-generative-adversarial-network-models/>

The GANs with Python EBook is where you'll find the **Really Good** stuff.



**Cristian** September 22, 2022 at 7:26 pm #

REPLY ↗

>> SEE WHAT'S INSIDE

It's fantastic! Thanks for this awesome tutorial!



**James Carmichael** September 23, 2022 at 5:53 am #

REPLY ↗

Thank you for the support and feedback Cristian! We greatly appreciate it!



**Calvin** October 18, 2022 at 6:10 pm #

REPLY ↗

Hello sir, thankyou for this amazing tutorial i really appreciate it.

But i have a question.....i have a task called face emotion generator using GAN. So basically it is a "generator" that can modified the emotion based on 1 static image input. For example, i have an angry

image. When i feed it into the “generator”, it can modified into another emotion such like sad, neutral, disgusting, etc.

### Never miss a tutorial:

disgusting, etc.

 thi GA the r method for my task?

Anw iam using AffectNet-HQ dataset from Kaggle.

### Picked for you:

 How to Develop a Pix2Pix GAN for

Image-to-Image Translation

**Calvin** October 18, 2022 at 10:42 pm #

REPLY ↗

Hi sir, thankyou for making this amazing tutorial.

How to Develop a 1D Generative

Adversarial Network From Scratch

I have a question sir i have a task to make a face emotion generator using GAN. Basically this generator can create an image face to display in another emotions. For example, from 1 static image of a face, it can be modified so that it displays sad, happy, or other emotions.

 How to Develop a CycleGAN for image-to-image translation

to-Image Translation with Keras

iam using the AffectNet-HQ dataset from kaggle....

Thanks before.

 How to Develop a Conditional GAN  
(cGAN) From Scratch

REPLY ↗

**James Carmichael** October 19, 2022 at 6:56 am #

How to Develop a Wasserstein

Generative Adversarial Network (WGAN)

Hi Calvin...You are very welcome! Yes, please proceed with CGAN and let us know what you find.

### Loving the Tutorials?

 The GANs with Python EBook is

where you'll find the **Really Good** stuff.

does the image result can be realistic if i only works with CGAN sir?

REPLY ↗

>> SEE WHAT'S INSIDE

### Leave a Reply

Name (required)

**Never miss a tutorial:**

Email (will not be published) (required)



SUBMIT COMMENT

**Picked for you:**

How to Develop a Pix2Pix GAN for Image-to-Image Translation With TensorFlow and Keras  
and I help developers get results with machine learning.  
[Read more](#)  
[How to Develop a 1D Generative Adversarial Network From Scratch in Keras](#)



How to Develop a CycleGAN for Image-to-Image Translation with Keras



How to Develop a Conditional GAN (cGAN) From Scratch



How to Develop a Wasserstein Generative Adversarial Network (WGAN) From Scratch

**Loving the Tutorials?**

The [GANs with Python](#) EBook is where you'll find the **Really Good** stuff.

[>> SEE WHAT'S INSIDE](#)

**Never miss a tutorial:****Picked for you:**

How to Develop a 1D Generative Adversarial Network From Scratch in Keras

© 2023 Guiding Tech Media. All Rights Reserved.

[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

How to Develop a CycleGAN for Image-  
to-Image Translation in Keras  
[Disclaimer](#) [Terms](#) [Contact](#) [Sitemap](#) [Search](#)



How to Develop a Wasserstein Generative Adversarial Network (WGAN)  
From Scratch

**Loving the Tutorials?**

The [GANs with Python](#) EBook is where you'll find the **Really Good** stuff.

[>> SEE WHAT'S INSIDE](#)